

# From Evidence to Reasoning: Trace-Augmented Reasoning under Conflicting Evidence

Anonymous ACL submission

## Abstract

Retrieval-augmented generation is unreliable when retrieved documents disagree, go stale, or fail to support any answer. We attribute this to a deficit in *judgment* rather than in answer generation, and separate the two empirically. We release a conflict-aware RAG benchmark of 1,679 examples and 10,343 retrieved documents, each annotated with a relevance judgment, an extracted supporting fact, a source-quality assessment, and complete annotator vote tallies. We supervise open-weight models to generate evidence traces comprising per-document adjudication, cross-document conflict analysis, and conflict-conditioned response synthesis, and evaluate them with CATS, which scores grounded refusal, behavioural adherence, factual grounding, and single-truth recall as separate components, each validated against an independent human study. Trace supervision transforms the refusal decision: refusal F1 rises from at most 0.429 to at least 0.826 across every model and prompt setting, and no baseline configuration exceeds the answer-everything baseline. Behaviour also becomes prompt-invariant: across a strict-to-minimal instruction ladder no trace-tuned model moves by more than the  $\pm 0.034$  measurement noise on any component, against baseline spreads of up to 0.823; grounding under the sparsest instruction is 0.792–0.857 tuned against 0.031–0.108 baseline. Answer quality does not improve: single-truth recall is unchanged ( $p = 0.77$ ), and the grounding advantage under well-specified prompts does not survive correction of two measurement artefacts we identify. Oracle ablations show that gold annotations for one stage move only the metrics that stage governs, and that the effects add, indicating the three stages are separable constructs. The human study finds behavioural adherence substantially less reliable than the other three components.

## 1 Introduction

Retrieval-augmented generation conditions a generator on a small set of retrieved snippets, ideally with

faithful attribution (Lewis et al., 2020; Borgeaud et al., 2022; Gao et al., 2024; Guu et al., 2020). In practice, retrieval rarely returns a coherent evidence set, instead producing documents that vary in relevance, quality, credibility, and recency. The retrieved documents may contain conflicting claims, provide only partial support, or collectively be insufficient to answer the user’s query.

Such a set demands a sequence of judgments before a single word of the answer is written: which document says what, how the documents relate to one another and whether what remains is sufficient. A system that treats the snippets as interchangeable context produces a fluent, confidently wrong answer; a system that refuses whenever sources disagree is no better, since disagreement is often compatible with a properly scoped answer.

Existing work addresses parts of this, but each stops short of a different part. *DRAGged into Conflicts* (Cattan et al., 2025) supervises the conflict label rather than the reasoning that produces it; *Chain-of-Note* (Yu et al., 2024) produces per-document notes without characterising how the retrieved set coheres; and *Trust-Score/Trust-Align* (Song et al., 2025) score refusal and citation behaviour without conditioning the expected response on the kind of conflict present. Closest in framing, Wu et al. (2026) also fine-tune open models on conflict signals, but derive those signals from generation probabilities rather than from an explicit textual record, Wang et al. (2025) resolve conflicting evidence by inference-time multi-agent debate rather than supervision, and Wang et al. (2026) pair a learned conflict detector with principled credibility scoring, where our source-quality field is a coarse URL heuristic. Muhamed et al. (2025) evaluate selective refusal under flawed context and independently find it to be a trainable capability, which is the deficit our results localise. The gap that remains is a single framework in which per-document reasoning, cross-document conflict in-

ference and justified abstention are supervised together and evaluated together.

We close that gap with a *structured evidence trace*. Before answering, the model exposes a three-stage record of its own adjudication (Figure 1): **Stage 1** assigns each snippet a verdict with an extracted key fact and a source-quality label; **Stage 2** aggregates those judgments into one of five conflict patterns, following the taxonomy of Cattán et al. (2025); **Stage 3** produces a citation-linked answer or a justified abstention, conditioned on that pattern (§3). The trace is emitted under a fixed output contract rather than as free-form reasoning, so it is machine-checkable and a failure is attributable to a stage rather than to the system as a whole. Liu et al. (2026) share this philosophy, training auditable traces by reinforcement learning on multi-hop QA rather than supervising conflict adjudication. Prompting can produce conflict-aware answers; it cannot produce conflict-aware judgment, and separating the two empirically is this paper’s central test.

**A benchmark that can carry the claim.** Testing this requires evaluation data the training distribution did not shape, so we built an independent 736-example benchmark with *zero* identifier overlap with training, verified on the released files rather than asserted. Because open-web retrieval cannot surface some conflict regimes, a curated track supplies them and its composition is reported per class (§4).

**Contributions.** (1) **A conflict-aware RAG resource** (§4): 943 training/validation examples and a disjoint 736-example benchmark, with 10,343 adjudicated evidence units carrying a verdict, key fact, snippet-anchored quote, rationale, source-quality label and the full annotator vote tally. (2) **Evidence that trace supervision buys judgment, not answer quality** (§8): refusal calibration and prompt-invariance improve; single-truth recall and grounding under well-specified prompts do not. (3) **CATS** (§6): an evaluation framework scoring a conflict-aware response as four components rather than one number, each reported over its own applicability denominator and separately validated against an independent human study. (4) **Oracle ablations** (§8.4): supplying gold information one stage at a time localises which stage limits the system and gives construct-validity evidence for the three-stage decomposition.




Figure 1: The evidence trace. The model adjudicates each retrieved document, infers the cross-document conflict pattern and produces a conflict-conditioned grounded answer or a justified refusal. Every stage is emitted in a machine-checkable form, so a downstream failure can be attributed to a stage rather than to the system as a whole.

## 2 Related Work

### 2.1 Knowledge and Evidence Conflicts

We scope this paper to conflicts *among retrieved documents*. Conflict between a model’s parametric knowledge and retrieved evidence, the context-memory setting targeted by much recent work (Huo et al., 2025; Zhang et al., 2025; Ye et al., 2026), is out of scope: our systems are evaluated only on whether they reason correctly over the supplied set.

Retrieval-augmented generation conditions a model on documents drawn from a large corpus at generation time (Guu et al., 2020; Lewis et al., 2020; Izacard et al., 2023; Borgeaud et al., 2022; Gao et al., 2024; Ram et al., 2023). Retrieved documents routinely carry conflicting information, which complicates generation and has become a research area of its own (Xu et al., 2024a). *ConflictBank* (Su et al., 2024) evaluates how models handle conflicts within retrieved knowledge, within parametric knowledge and in their interplay, pairing Wikidata claims with naturally occurring or semantically constructed contradictory evidence. *WikiContradict* (Hou et al., 2024) benchmarks factual and temporal robustness on real-world contradictions between Wikipedia passages, with human-annotated instances. *ECON* (Jiayang et al., 2024) studies detection and resolution of evidence conflicts directly. *DRAGged into Conflicts* (Cattán et al., 2025) is closest to our setting: it proposes a conflict taxonomy spanning temporal, debatable, misinformation and complementary cases and demonstrates that models often fail to adapt their responses without explicit conflict-aware guidance.

Most of this literature is concerned with *detecting* or *classifying* conflict. Our framework instead supervises the reasoning that produces the classification and, because that reasoning is emitted in a checkable format, evaluates whether a model’s document-level judgments, its conflict inference and its final response behaviour are mutually consistent, rather than scoring only the final answer.

## 2.2 Reasoning Supervision, Multi-Hop and Robustness in RAG

Several lines of work supervise or elicit intermediate reasoning. *Chain-of-Thought* prompting established that generating intermediate steps improves complex reasoning (Wei et al., 2022). *Chain-of-Note* (Yu et al., 2024) adapts this to retrieval by generating sequential reading notes for each retrieved document, improving robustness to noisy and irrelevant passages and enabling a model to answer “unknown” when neither retrieved nor intrinsic knowledge suffices. Adaptive adversarial training (Fang et al., 2024) and robustness training (Yoran et al., 2024) mitigate retrieval noise and *GSM-IC* (Shi et al., 2023) shows that irrelevant context substantially degrades reasoning.

Our work differs in two respects. First, each stage has a closed label set, so Stage-1 and Stage-2 decisions are separately scorable against gold annotations rather than merely inspectable. Second and more importantly, *we treat the reasoning trace as an object of evaluation rather than a means of improving the answer*. The three-stage decomposition earns its place only if the stages are genuinely separable, which we test directly by supplying gold information for one stage at a time and asking whether each intervention moves the metrics that stage should govern and no others (§8.4).

## 2.3 Measuring Trust, Groundedness and Conflict Alignment

Trustworthiness metrics for RAG largely score an answer against the retrieved set as a whole; what conflict-aware evaluation additionally requires is that the *expected* response depend on the kind of disagreement present, so that a balanced summary and a decisive correction are not scored by the same rule.

Evaluations of RAG frequently assess end-to-end system quality (Gao et al., 2024; Xu et al., 2024b), conflating retriever quality with generator capability (Fan et al., 2024). *Trust-Score* (Song et al., 2025) advances this by evaluating trustwor-

thiness holistically through answer correctness, citation reliability and refusal behaviour and its companion *Trust-Align* aligns models to improve that score; its held-out refusal material forms the refusal track of our resource. *FaithEval* (Ming et al., 2025) probes whether models stay faithful to context under adversarial or counterfactual conditions.

CATS extends this line by conditioning the expected response on the inferred conflict type, evaluating grounding at claim level with citation sensitivity, and reporting every component over its own applicability denominator, since the four constructs apply to different subsets of examples and conflating them confounds the comparison the framework is meant to support (§8.3). We validate the judges against an independent human study (Zheng et al., 2023) and report which components that validation supports and which it does not (§9).

## 3 Task Formulation and the Evidence Trace

An example consists of a query  $q$  and a retrieved set  $D = \{d_1, \dots, d_N\}$ . The system must produce either a citation-linked answer or the controlled refusal string CANNOT ANSWER, INSUFFICIENT EVIDENCE, reasoning only over the retrieved evidence and refusing when that evidence is insufficient. Rather than mapping  $(q, D)$  directly to a response, we supervise an intermediate representation emitted as part of the output.

**Stage 1: document adjudication.** Every  $d_i$  receives a verdict in  $\{\text{supports}, \text{partially supports}, \text{irrelevant}\}$ , a one-sentence key fact entailed by the snippet, and a coarse source-quality label. The *partially supports* class is deliberately broad, covering on-topic but incomplete, hedged, scoped or single-facet evidence, because collapsing retrieval into a binary useful/useless distinction produces the failure mode we care about: refusing whenever no single snippet is individually decisive.

**Stage 2: conflict analysis.** The Stage-1 judgments are aggregated into exactly one of five labels, each prescribing a different aggregation strategy; Table 11 (Appendix F) gives each type’s definition and the response behaviour it requires. The taxonomy separates *complementarity* from *conflict*: two snippets can contribute different necessary pieces without disagreeing.




Figure 2: System overview. A weighted multi-annotator committee produces the supervision corpus; models are fine-tuned to emit the three-stage trace; CATS evaluates refusal calibration, conflict behaviour, factual grounding and answer recall, with an independent human study validating the automated judgements.

**Stage 3: grounded synthesis.** Stages 1 and 2 are emitted *inside* the `<think>` block; Stage 3 is the response itself and sits outside it. The final response is conditioned on the conflict label, with sentence-level `[dX]` citations, or is the exact refusal string, terminating with `[ [END-OF-ANSWER] ]` (Figure 2 places this in the overall pipeline).

## 4 Dataset Generation and Validation

This section describes the generated dataset and how it was validated (Table 1). Appendix A gives the complete construction record, per-source distributions, retrieval protocol and human-review statistics.

### 4.1 Training and validation

The 943-example training release (862 train / 81 validation) unions three source families: 458 normalised examples from the CONFLICTS collection (Cattan et al., 2025), 272 held-out refusal examples derived from Trust-Align (Song et al., 2025) and 213 examples drawn from the benchmark construction stream that were *not* admitted to the final benchmark. All 213 were checked for identifier overlap against the released benchmark; the intersection is empty.

### 4.2 An independent benchmark

The 736-example benchmark is built in three tracks and reporting them separately is deliberate.

| Conflict type         | Train      |      | Benchmark  |      |
|-----------------------|------------|------|------------|------|
|                       | <i>n</i>   | %    | <i>n</i>   | %    |
| No conflict           | 343        | 36.4 | 211        | 28.7 |
| Complementary         | 286        | 30.3 | 221        | 30.0 |
| Conflicting opinions  | 158        | 16.8 | 109        | 14.8 |
| Outdated              | 136        | 14.4 | 158        | 21.5 |
| Misinformation        | 20         | 2.1  | 37         | 5.0  |
| Answerable            | 649        | 68.8 | 608        | 82.6 |
| Refusal-required      | 294        | 31.2 | 128        | 17.4 |
| <b>Records</b>        | <b>943</b> |      | <b>736</b> |      |
| Retrieved documents   | 6,642      |      | 3,701      |      |
| Docs / record (mean)  | 7.04       |      | 5.03       |      |
| Unique source domains | 2,633      |      | 1,431      |      |

Table 1: **Dataset composition and scale.** Conflict-type and answerability distribution of the training split (943 records, including validation) and the held-out benchmark (736). The benchmark is deliberately harder than the training distribution: fewer no-conflict cases (28.7% against 36.4%), a  $1.5\times$  larger outdated-information share and a  $2.4\times$  larger misinformation share. Full per-source composition in Appendix A.

**Track A: web-retrieved (508).** Queries are drawn from ConflictingQA (Wan et al., 2024), SituatedQA (Zhang and Choi, 2021), FreshQA (Vu et al., 2024) and QACC (Liu et al., 2025), with every query overlapping the training pool excluded. Documents are retrieved live and a query-focused window is selected by a TAS-B dense bi-encoder (Hofstätter et al., 2021). Appendix A gives the full funnel and the five-document sampling rule.

**Track B: refusal (128).** Refusal-required examples come from the held-out refusal pool of Trust-Align (Song et al., 2025) and are *not* web-retrieved.

**Track C: curated conflict (100).** Open-web retrieval does not reliably surface every conflict regime, one cannot query a search index for the version of a fact that was true two years ago. We therefore add 78 examples built from dated Wikipedia revision pairs, in the spirit of Hou et al. (2024), 19 multi-hop contexts from HotpotQA (Yang et al., 2018) and three medical/misinformation cases. The track is strongly class-targeted: all 78 revision-pair examples are *outdated-information* cases, supplying 49% of that class and all 19 multi-hop examples are *complementary* (Table 5, Appendix A).

### 4.3 Annotation depth

Each of the 10,343 retrieved documents carries six fields: an adjudicated *verdict*, an extracted *key fact*, a snippet-anchored *supporting quote*, a *rationale*,



a *source-quality label* and the complete *annotator vote tally*. Labels come from a weighted multi-annotator committee that votes on the categorical decision at each stage and then adopts the entire explanation bundle from the highest-weight member on the winning side, so a verdict, its quote and its rationale always come from one coherent interpretation rather than being stitched across annotators. The annotation contract requires the key fact to be entailed by its quote and the quote to be a contiguous verbatim span of at most fifty words (Appendix B); the prompts that enforce it are released with the resource (Appendix M).

#### 4.4 Human review

Two human workflows, kept separate, cover the two populations. Seven reviewers screened the benchmark candidates, recording a retention decision, a preliminary conflict type and eight quality judgments: retrieval quality, evidence sufficiency, conflict clarity, query specificity, source reliability, the count of relevant documents, gold-answer feasibility and reviewer confidence. Five-way conflict-type agreement is 94.77% ( $\kappa = 0.92$ ) and retention-decision agreement is 98.21% ( $\kappa = 0.92$ ). Training-side review assigned each base example to two reviewers, giving conflict-type agreement of 83.46% ( $\kappa = 0.77$ ). The benchmark second pass is *sequential*, so its figures measure verification consistency rather than blind double annotation. Appendix C gives the full protocol and per-field agreement tables.

### 5 Trace-Supervised Fine-Tuning

#### 5.1 Objective and prompt variants

The supervision target is the serialised trace followed by the final response. Three prompt variants are used: **strict** states the full taxonomy, verdict heuristics and output contract (1,965 words); **runtime** keeps a compact structural schema (195 words); **minimal** supplies only the query, the documents and a four-sentence evidence policy (40 words). The mixture also includes direct document-verdict, conflict-type and final-response supervision as separate task views, so each component decision receives explicit signal in isolation as well as inside the full trace. We train *and* evaluate across all three variants, which is what makes it testable whether conflict-aware behaviour is *learned* or merely *recited* from the instruction; §8.2 reports the difference, and Appendix N gives the three prompts

verbatim together with the design rationale.

#### 5.2 Fine-tuning

All models are fine-tuned with QLoRA (Detrmers et al., 2023; Hu et al., 2021) on the standard token-level cross-entropy of the trace and response, with no auxiliary loss. Appendix D gives the hyperparameters, the checkpoint selection criterion and the two training mixtures (recipes K and L), which differ by model family and are needed to interpret the per-model results in §8.

### 6 CATS: Conflict-Aware Trust Score

Standard RAG metrics collapse distinct failures: answering when evidence is absent, mishandling the conflict type, making unsupported claims and omitting the target answer are all “a bad answer”. CATS separates them into four components with independent applicability.

---

#### Grounded Refusal (GR)

whether the system answered exactly when the evidence permits answering; scored deterministically from the answerability label

#### Behaviour Adherence (BA)

whether the answer followed the response policy its conflict type requires

#### Factual Grounding (FG)

whether each extracted claim is both supported by an eligible document *and* cited

#### Single-Truth Recall (STR)

whether the response adopts the target answer as its own conclusion, rather than quoting or attributing it

---

Table 2: **The four CATS components**, each reported over its own applicability denominator because the constructs apply to different subsets of examples. Appendix F gives the full definitions, FG eligibility rules, per-type BA policies and the committee voting rule.

Table 2 states what each component measures and the population it is scored over. The definitions are operational: FG counts a claim only when the supporting document’s *gold* verdict is *supports* or *partially supports*, BA is judged against a different expected response for each conflict type, and STR excludes conflicting-opinion items because no single conclusion exists for them. Appendix F sets out each rule; the denominators differ, so any comparison against another framework should start there.

**Judges:** BA, FG and STR are judged by a committee of three locally served open-weight models with fixed weights, aggregated by weighted majority; factual-grounding support additionally requires

corroboration by at least two judges and GR is never committee-judged. We compare this against human review in §9. The complete judge bundle is one system message plus three task prompts and the five-type behaviour rubric, reproduced verbatim in Appendix O; metric definitions and aggregation are in Appendix F.

**Aggregate scores:** the four components can also be combined into a single number per example, gated by the answer/refuse decision so that a wrong decision scores zero. We define two averages, **CATS-Prevalence** and **CATS-Balanced**, in Appendix F.8, together with the construction and its structural properties. We do not report either in the main results: because a refusal-required example is scored on its decision alone, such examples carry 40% of CATS-Balanced while making up 17.4% of the benchmark, so a model can raise the aggregate by improving refusals without improving any answer. Every claim in §8 is stated on the components.

## 7 Experimental Setup

### 7.1 Primary evaluation

We fine-tune four instruction-tuned open-weight models, Qwen2.5-7B and Qwen2.5-32B (Yang et al., 2024), Llama-3.1-8B (Llama Team, AI at Meta, 2024) and Mistral-7B (Jiang et al., 2023), and compare each against its baseline version on the four CATS components. Within a cell the benchmark examples, prompts, decoding policy, retry policy, post-processing and evaluator are identical; the only difference is the presence of the LoRA adapter, so any difference is attributable to trace supervision. Each pair is evaluated under the strict and minimal prompts, with the intermediate runtime prompt reported in Appendix Q.

Generation is deterministic (temperature 0) with profile-specific continuation budgets and sentinel-based stopping; retries only enlarge the token budget for the identical prompt, never resampling or repairing content (Appendix E). Per-example scores lie in  $[0, 1]$ , so a distribution-free 95% interval half-width at  $n = 736$  is  $\pm 0.034$ ; we treat differences below this as uninformative, report exact sign tests for paired comparisons across the twelve model  $\times$  prompt cells, and use exact Clopper–Pearson intervals on the 128-example refusal subgroup.

### 7.2 Ablation studies

**Oracle ablations.** To test whether the three stages are separable, we re-run the fine-tuned models with gold information injected one stage at a time: **+conflict label** supplies the gold Stage-2 type, **+per-doc notes** the gold Stage-1 verdicts and key facts, and **+both** supplies each. If the stages are distinct, an intervention should move the metrics its stage governs and leave the rest unchanged (§8.4).

**Comparison with prompting methods.** Two prompting-only strategies on baseline models separate what supervision provides from what an elicitation prompt can. *Chain-of-Note* (Yu et al., 2024) writes per-document reading notes before answering, the closest analogue to our Stage 1; *few-shot Chain-of-Thought* (Wei et al., 2022) supplies worked exemplars of the reasoning format. Both receive the identical benchmark, decoding policy and evaluator.

## 8 Results

Table 3 reports the end-to-end matrix, each metric as a base/SFT column pair; Appendix Q gives all twelve model  $\times$  prompt cells.

| Prompt              | Refusal F1 |              | BA    |              | FG    |              | STR   |              |
|---------------------|------------|--------------|-------|--------------|-------|--------------|-------|--------------|
|                     | base       | SFT          | base  | SFT          | base  | SFT          | base  | SFT          |
| <b>Qwen2.5-7B</b>   |            |              |       |              |       |              |       |              |
| strict              | 0.429      | <b>0.911</b> | 0.811 | 0.780        | 0.617 | <b>0.838</b> | 0.561 | <b>0.687</b> |
| minimal             | 0.031      | <b>0.881</b> | 0.737 | <b>0.768</b> | 0.031 | <b>0.837</b> | 0.757 | 0.686        |
| <b>Qwen2.5-32B</b>  |            |              |       |              |       |              |       |              |
| strict              | 0.367      | <b>0.905</b> | 0.789 | <b>0.821</b> | 0.851 | 0.812        | 0.733 | <b>0.749</b> |
| minimal             | 0.015      | <b>0.911</b> | 0.747 | <b>0.809</b> | 0.108 | <b>0.822</b> | 0.773 | 0.738        |
| <b>Llama-3.1-8B</b> |            |              |       |              |       |              |       |              |
| strict              | 0.255      | <b>0.926</b> | 0.721 | <b>0.805</b> | 0.705 | <b>0.852</b> | 0.634 | <b>0.727</b> |
| minimal             | 0.015      | <b>0.934</b> | 0.743 | <b>0.795</b> | 0.074 | <b>0.857</b> | 0.710 | <b>0.714</b> |
| <b>Mistral-7B</b>   |            |              |       |              |       |              |       |              |
| strict              | 0.352      | <b>0.826</b> | 0.660 | <b>0.793</b> | 0.624 | <b>0.796</b> | 0.563 | <b>0.644</b> |
| minimal             | 0.000      | <b>0.839</b> | 0.727 | <b>0.786</b> | 0.079 | <b>0.792</b> | 0.700 | 0.627        |

Table 3: **End-to-end results on the 736-example held-out benchmark.** Within each pair the benchmark, prompts, decoding policy and evaluator are identical; only the LoRA adapter differs. Bold marks a fine-tuned value that exceeds its baseline. Differences below the  $\pm 0.034$  noise floor (§7) are uninformative. We show the two endpoints of the instruction ladder; the intermediate runtime condition is in Appendix Q.

### 8.1 Refusal recall improves substantially

The largest and cleanest effect is the answer/refuse decision. Refusal F1 separates completely: every baseline configuration scores  $\leq 0.429$  and every

trace-tuned configuration  $\geq 0.826$ , with *no overlap*, a mean improvement of  $+0.709$ . This is not over-caution. Decision accuracy also rises in every cell and crosses the answer-everything floor: the benchmark is 82.6% answerable, so answering everything scores 0.826 for free, and no baseline configuration exceeds that floor by more than the  $\pm 0.034$  noise interval while every trace-tuned cell reaches 0.927–0.976. Behaviour Adherence improves alongside the decision but unevenly, from  $+0.096$  for Mistral-7B to  $+0.000$  for Qwen2.5-7B, an ordering that tracks baseline instruction-following. Appendix Q gives the per-cell recall and accuracy figures and the Llama-3.1-8B case that shows why the two decision numbers must be read together.

## 8.2 Learnt behaviour is prompt-invariant

Removing the prompt scaffolding is close to a null operation for fine-tuned models and close to catastrophic for baseline ones. The minimal condition is a genuinely sparse instruction, forty words that never name the trace, the taxonomy or the citation format, so a model that still produces them is reproducing a learned mapping rather than following the prompt. Component by component, no trace-tuned model varies by more than the  $\pm 0.034$  noise floor across the ladder, against baseline spreads of 0.240–0.398 on refusal F1 and 0.620–0.823 on factual grounding; under the minimal prompt baseline grounding falls to 0.031–0.108 while trace-tuned models sustain 0.792–0.857 (Table 18). The result also reproduces on citation measurements built after and independently of CATS. Figure 13 plots the full ladder, and Appendix J gives the runtime-condition and citation-sufficiency figures.

## 8.3 What trace supervision does *not* buy

Two components show no consistent gain. Single-truth recall is positive in only 7 of 12 cells (sign test  $p = 0.77$ ), a well-powered null: fine-tuned models are not recovering target answers more often. Factual grounding is not directly comparable across run types, because it is averaged over  $fg\_n = 608 + \text{missed refusals}$ , so a baseline model that rarely refuses is additionally scored on  $\sim 130$  items whose evidence is insufficient by construction; bounding FG on the like-for-like population, the intervals overlap in 6 of 12 cells, every one of them a strict or runtime cell (Table 17; see Appendix I). Trace supervision therefore raises grounding when the instruction is sparse, but under

a well-specified prompt no consistent advantage is established. Given a well-specified prompt a baseline model already produces competitive grounded answers; what it lacks is the judgment to decide whether to answer at all, and the ability to retain any of this without being told the procedure. That is what trace supervision supplies, which is why these negative results are predictions of our account rather than exceptions to it. Appendix I gives the per-cell bounds and the instruction-volume effect for baseline models.

## 8.4 Mechanism-matched oracle ablations

This ablation asks what the final response gains when the model is handed the gold conflict type, the gold per-document notes, or both, instead of having to infer them. Each condition replaces one stage of the trace with ground truth and leaves the rest of the pipeline unchanged, so the resulting change isolates what that stage contributes. They are diagnostic interventions, never deployment settings.

Read as a single aggregate, their effect is small and the analysis stops there. Read *per metric*, with metrics the intervention should not affect acting as a specificity control, the result is considerably stronger. Table 19 and Figure 17 give the per-metric effect of each oracle, on which the paragraphs below draw; Appendix K discusses them.

**The conflict label moves behaviour and only behaviour.** Supplying the gold Stage-2 label improves Behaviour Adherence ( $+0.013$ , 8/9 cells,  $p = 0.039$ ) and moves no other metric significantly; grounding, citation and decision metrics are flat to slightly negative.

**The document notes move grounding and only grounding.** Supplying gold Stage-1 notes improves Factual Grounding ( $+0.047$ ), answer quality ( $+0.044$ ), strict citation quality ( $+0.041$ ), citation precision ( $+0.033$ ) and citation sufficiency ( $+0.028$ ), each positive in *every* cell. Behaviour shows a small spillover, expected since correct document verdicts also constrain the conflict inference and the answer/refuse decision is untouched, as it should be since GR is upstream of both interventions.

**Interpretation.** That each intervention moves the metrics its stage governs and no others, and that the effects add ( $|\Delta_{\text{both}} - (\Delta_{\text{conflict}} + \Delta_{\text{notes}})| \leq 0.012$ ), is a construct-validity argument for the decompo-

sition. For *baseline* models every oracle instead *hurts*, catastrophically so for grounding under the conflict label ( $-0.238$ ), a second internalization result. Appendix K gives the per-metric residuals, the behaviour exception, the caveat on where the remaining error lives, and the baseline breakdown.

## 8.5 External prompting comparators

We compare against two prompting-only strategies applied to baseline models: few-shot Chain-of-Thought (Wei et al., 2022) and Chain-of-Note-style retrieval notes (Yu et al., 2024). Neither is fine-tuned; both receive the same benchmark and evaluator (Table 4).

| Method              | Ref. F1      | BA           | FG           | STR          |
|---------------------|--------------|--------------|--------------|--------------|
| <i>Llama-3.1-8B</i> |              |              |              |              |
| CoN (prompt-only)   | 0.000        | 0.632        | 0.609        | 0.658        |
| CoT (few-shot)      | 0.302        | 0.716        | 0.635        | 0.678        |
| TRACE-SFT (ours)    | <b>0.926</b> | <b>0.805</b> | <b>0.852</b> | <b>0.727</b> |
| <i>Mistral-7B</i>   |              |              |              |              |
| CoN (prompt-only)   | 0.000        | 0.685        | 0.522        | 0.676        |
| CoT (few-shot)      | 0.525        | 0.657        | 0.417        | 0.528        |
| TRACE-SFT (ours)    | <b>0.826</b> | <b>0.793</b> | <b>0.796</b> | <b>0.644</b> |
| <i>Qwen2.5-7B</i>   |              |              |              |              |
| CoN (prompt-only)   | 0.000        | 0.727        | 0.666        | 0.719        |
| CoT (few-shot)      | 0.402        | <b>0.781</b> | 0.734        | <b>0.724</b> |
| TRACE-SFT (ours)    | <b>0.911</b> | 0.780        | <b>0.838</b> | 0.687        |

Table 4: **External prompting comparators** on baseline models; TRACE-SFT rows use the strict prompt, so every row receives a well-specified instruction. Both comparators reach usable answer content but neither acquires the answer/refuse decision.

Both reach usable answer content, Chain-of-Note reaching factual grounding of 0.52–0.67, but neither acquires the answer/refuse decision. Chain-of-Note *never abstains*: zero correct refusals out of 128 on all three models, landing exactly on the answer-everything floor. Few-shot Chain-of-Thought abstains inconsistently rather than rarely, from 32/128 to 85/128 across models. Structured prompting can therefore elicit conflict-aware *answers* without supplying conflict-aware *judgment*, the distinction §1 draws. Appendix H gives the per-model abstention spread, why the Mistral-7B exception is a refusal prior rather than acquired judgment, and Chain-of-Note’s citation hygiene, the poorest of any method examined.

**Citation behaviour.** Trace supervision changes citation *sufficiency*, not citation precision: precision is near-ceiling for every model regardless of tuning, while the recall side moves sharply. Appendix H

gives the seven-metric analysis, the definitions and the per-cell values.

## 9 Validating the Evaluator

A judged metric is only as trustworthy as its judge, so we validated the CATS committee against an independent human study: four reviewers double-judged 300 responses spanning both model families, both training conditions and all five conflict types. The resulting reliability ordering is identical on two independent instruments and bounds what we may claim. **Factual Grounding and Single-Truth Recall are reliable** ( $\kappa = 0.78$  and  $0.79$  against humans); **Behaviour Adherence is not** ( $\kappa = 0.41$ ), with the weakness concentrated in Complementary information and Misinformation, where human reviewers also disagree most. We therefore report conflict-conditioned behaviour qualitatively and do not interpret individual BA margins. Appendices G and P give the study design, every agreement statistic and the coverage caveats; any reading of a BA number should be checked against them.

## 10 Conclusion

Trace supervision transforms the answer/refuse decision and makes conflict-aware behaviour survive removal of the instruction describing it; it does *not* improve answer writing, because that deficit was never there to fix; and the three stages behave as separable constructs, so the trace is a measurement instrument rather than a formatting convention.

Three questions remain open. The residual error after both oracles is small but unattributed, and separating Stage-3 synthesis failure from annotation noise and judge strictness needs a targeted audit. Behaviour Adherence is the weakest instrument we report, so conflict-behaviour claims should stay qualitative until it improves. And misinformation remains the smallest class in the benchmark; expanding it without perturbation-based synthesis is an open data-collection problem.

The practical consequence is a separation of concerns. A retrieval system that answers well but cannot decide *when* to answer is not a weaker version of one that can; the two differ in a capability that no prompt in our ladder supplied, and that supervising the reasoning did. Building such a system means supervising the judgment, and evaluating it means scoring that judgment apart from the answer it produces.



## Limitations

**Per-document verdicts are never human-validated.** The Stage-1 verdicts (*supports*, *partially supports*, *irrelevant*) that gate Factual Grounding eligibility are produced entirely by the annotation committee, including frontier-model annotators, and are checked by human reviewers at no stage. This is a different status from the conflict-type label, which is reviewed twice, and it means the eligibility set on which every grounding and citation number depends rests on model agreement rather than human adjudication. The refusal-track validity check (Appendix C.3) is indirect evidence that the verdicts track something real, not a substitute for direct validation.

**Evidence-bounded evaluation.** CATS scores whether an answer is justified by the retrieved set, not whether it is true. A globally correct answer unsupported by  $D$  scores poorly. This is deliberate for retrieval-faithfulness research and inappropriate if the question is unrestricted factual competence. Within that scope, grounding is measured permissively: the judge is given seven routes to support and instructed to bias toward inclusion (Appendix O), so our grounding numbers are an upper bound relative to a strict entailment criterion.

**Behaviour Adherence reliability.** As §9 reports, BA is the least reliable component ( $\kappa = 0.41$  against humans,  $\alpha = 0.46$  among judges) and all observed BA differences fall inside that disagreement envelope. BA-based conclusions here are qualitative. Relatedly, the judge committee uses a priority-weighted rule under which the highest-priority judge cannot be overruled by the other two on behaviour judgments once its self-reported confidence exceeds 5/6; behaviour is therefore closer to a single-judge measurement than the panel size suggests, precisely in the disputed cases.

**Shared model families between annotation and evaluation.** The annotation committee and the evaluation committee draw on the same three open-weight model families, with different weightings. The constructs differ (document adjudication versus answer judging) and Grounded Refusal is deterministic and outside the committee entirely, but the evaluator is partly scoring labels produced by related models. The human study in §9 is our external check on this and it covers two of the four model families.

**Aggregate properties.** The type-balanced aggregate assigns 40% of its weight to refusal-required examples, which score decision correctness only, against 17.4% of the data; one type-3 refusal example carries roughly twelve times the influence of one type-1 answerable example; and the misinformation class, 5% of the data, determines 20% of the score. A model with zero usable answer content but perfect refusal calibration would score 0.400. This is why we report components first and aggregates second (Appendix F.8).

**Recipe asymmetry and scale.** Qwen models use recipe K and Llama/Mistral recipe L; no completed Qwen matrix exists for recipe L. Recipes J/K/L change several ingredients together and are principled recipe revisions rather than single-factor ablations. All models are in the 7B–32B range and Mistral-7B remains over-abstaining under every recipe we tried, a model-family limitation we report rather than omit.

**Annotation quality.** Supporting quotes are snippet-*anchored* rather than fully verbatim: 81–83% are exact contiguous substrings and ~96% are traceable to the snippet, with ~4% showing no detectable overlap. Source-quality labels come from an explicit URL whitelist rule (Appendix B), not from an independent credibility assessment, and are systematically pushed toward *low* on the refusal track, whose upstream source uses placeholder URLs. Three benchmark records carry a numeric conflict identifier inconsistent with their authoritative textual label; we verified this did not affect any reported result.

**Human review coverage.** The benchmark second review pass is sequential rather than blind, so its agreement figures measure verification consistency. Of the 943 consolidated training rows, 285 are consensus-completed rather than genuine two-reviewer pairs and 105 of the 608 answerable benchmark records appear only as consensus-completed entries. Several funnel stages (1,878 assigned candidates against 1,221 reviewed; 200 curated refusals against 128 released) have no retained per-record selector; we report the counts as observed rather than inferring a rule.

**Comparator scope.** The prompting comparators cover three of the four model families and one prompt configuration each; they are not exhaustive prompt-engineering baselines. One comparator metric shifted substantially between our earlier

evaluation and this one in a way most consistent with a corrected answer-extraction step, so we base the comparator conclusion on abstention behaviour and citation hygiene, which are stable across both, rather than on grounding alone.

## Ethical considerations

**Data provenance and licensing.** The released resource is derived from publicly available research datasets and from live web retrieval. Source provenance is preserved in every record and the dataset cards identify the upstream collection for each source family so that downstream users can consult the relevant licences before redistributing source-derived material. Retrieved snippets are stored with their source URL and, where available, a timestamp; because live retrieval is time-varying, the release should be treated as a snapshot rather than a reproducible query against a current index.

**Human annotators.** Human review was carried out by members of the research group rather than by crowdworkers. Reviewer identifiers are retained in the review artifacts for auditability but carry no demographic information and we make no claims about recruitment, compensation or demographics that our records do not support.

**Intended use and risks.** The resource is intended for research on evidence-grounded generation, abstention calibration and conflict-aware evaluation. Two cautions apply. Conflict labels and answerability judgments describe the *supplied* evidence, not world truth and must not be read as factual adjudications of the underlying questions, several of which concern contested scientific or medical topics. Second, the source-quality field is a coarse URL-whitelist heuristic retained as annotation metadata; it is not a credibility ranking of publishers and should not be used as one.

**Disclosure of AI assistance.** Frontier language models were used to generate the synthetic annotations in our dataset and to evaluate model performance as LLM-as-a-judge; the methodology for both is specified in §4 and §6. They were also used for grammatical refinement of the text. All experimental design, analysis and claims are the authors’ own and the paper is an accurate representation of the underlying work.

## References

- Sebastian Borgeaud, Arthur Mensch, Jordan Hoffmann, Trevor Cai, Eliza Rutherford, Katie Millican, George van den Driessche, Jean-Baptiste Lespiau, Bogdan Damoc, Aidan Clark, Diego de Las Casas, Aurelia Guy, Jacob Menick, Roman Ring, Tom Hennigan, Saffron Huang, Loren Maggiore, Chris Jones, Albin Cassirer, and 9 others. 2022. [Improving language models by retrieving from trillions of tokens](#). *Preprint*, arXiv:2112.04426.
- Arie Cattan, Alon Jacovi, Ori Ram, Jonathan Herzig, Roei Aharoni, Sasha Goldshtein, Eran Ofek, Idan Szpektor, and Avi Caciularu. 2025. [DRAGed into conflicts: Detecting and addressing conflicting sources in search-augmented LLMs](#). *Preprint*, arXiv:2506.08500.
- Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. 2023. [QLoRA: Efficient finetuning of quantized LLMs](#). *Preprint*, arXiv:2305.14314.
- Wenqi Fan, Yujuan Ding, Liangbo Ning, Shijie Wang, Hengyun Li, Dawei Yin, Tat-Seng Chua, and Qing Li. 2024. [A survey on RAG meeting LLMs: Towards retrieval-augmented large language models](#). *Preprint*, arXiv:2405.06211.
- Feiteng Fang, Yuelin Bai, Shiwen Ni, Min Yang, Xiaojun Chen, and Ruifeng Xu. 2024. [Enhancing noise robustness of retrieval-augmented language models with adaptive adversarial training](#). *Preprint*, arXiv:2405.20978.
- Yunfan Gao, Yun Xiong, Xinyu Gao, Kangxiang Jia, Jinliu Pan, Yuxi Bi, Yi Dai, Jiawei Sun, Meng Wang, and Haofen Wang. 2024. [Retrieval-augmented generation for large language models: A survey](#). *Preprint*, arXiv:2312.10997.
- Kelvin Guu, Kenton Lee, Zora Tung, Panupong Pasupat, and Ming-Wei Chang. 2020. REALM: Retrieval-augmented language model pre-training. In *Proceedings of the 37th International Conference on Machine Learning*.
- Sebastian Hofstätter, Sheng-Chieh Lin, Jheng-Hong Yang, Jimmy Lin, and Allan Hanbury. 2021. Efficiently teaching an effective dense retriever with balanced topic aware sampling. In *Proceedings of the 44th International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 113–122.
- Yufang Hou, Alessandra Pascale, Javier Carnerero-Cano, Tigran Tchakian, Radu Marinescu, Elizabeth Daly, Inkit Padhi, and Prasanna Sattigeri. 2024. WikiContradict: A benchmark for evaluating LLMs on real-world knowledge conflicts from Wikipedia. In *Advances in Neural Information Processing Systems 37: Datasets and Benchmarks Track*, pages 109701–109747.
- Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and

|     |                                                                                                                                                                                                                                                                                                                                                                                                             |     |
|-----|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----|
| 862 | Weizhu Chen. 2021. <a href="#">LoRA: Low-rank adaptation of large language models</a> . <i>Preprint</i> , arXiv:2106.09685.                                                                                                                                                                                                                                                                                 | 919 |
| 863 |                                                                                                                                                                                                                                                                                                                                                                                                             | 920 |
| 864 | Nan Huo, Jinyang Li, Bowen Qin, Ge Qu, Xiaolong Li, Xiaodong Li, Chenhao Ma, and Reynold Cheng. 2025. <a href="#">Micro-Act: Mitigate knowledge conflict in question answering via actionable self-reasoning</a> . In <i>Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)</i> , Vienna, Austria. Association for Computational Linguistics.  | 921 |
| 865 |                                                                                                                                                                                                                                                                                                                                                                                                             | 922 |
| 866 |                                                                                                                                                                                                                                                                                                                                                                                                             | 923 |
| 867 |                                                                                                                                                                                                                                                                                                                                                                                                             | 924 |
| 868 |                                                                                                                                                                                                                                                                                                                                                                                                             | 925 |
| 869 |                                                                                                                                                                                                                                                                                                                                                                                                             | 926 |
| 870 |                                                                                                                                                                                                                                                                                                                                                                                                             | 927 |
| 871 |                                                                                                                                                                                                                                                                                                                                                                                                             | 928 |
| 872 | Gautier Izacard, Patrick Lewis, Maria Lomeli, Lucas Hosseini, Fabio Petroni, Timo Schick, Jane Dwivedi-Yu, Armand Joulin, Sebastian Riedel, and Edouard Grave. 2023. Atlas: Few-shot learning with retrieval augmented language models. <i>Journal of Machine Learning Research</i> , 24(251):1–43.                                                                                                         | 929 |
| 873 |                                                                                                                                                                                                                                                                                                                                                                                                             | 930 |
| 874 |                                                                                                                                                                                                                                                                                                                                                                                                             | 931 |
| 875 |                                                                                                                                                                                                                                                                                                                                                                                                             | 932 |
| 876 |                                                                                                                                                                                                                                                                                                                                                                                                             | 933 |
| 877 |                                                                                                                                                                                                                                                                                                                                                                                                             | 934 |
| 878 | Neel Jain, Ping yeh Chiang, Yuxin Wen, John Kirchenbauer, Hong-Min Chu, Gowthami Somepalli, Brian R. Bartoldson, Bhavya Kailkhura, Avi Schwarzschild, Aniruddha Saha, Micah Goldblum, Jonas Geiping, and Tom Goldstein. 2023. <a href="#">NEFTune: Noisy embeddings improve instruction finetuning</a> . <i>Preprint</i> , arXiv:2310.05914.                                                                | 935 |
| 879 |                                                                                                                                                                                                                                                                                                                                                                                                             | 936 |
| 880 |                                                                                                                                                                                                                                                                                                                                                                                                             | 937 |
| 881 |                                                                                                                                                                                                                                                                                                                                                                                                             | 938 |
| 882 |                                                                                                                                                                                                                                                                                                                                                                                                             | 939 |
| 883 |                                                                                                                                                                                                                                                                                                                                                                                                             | 940 |
| 884 |                                                                                                                                                                                                                                                                                                                                                                                                             | 941 |
| 885 | Albert Q. Jiang, Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Florian Bressand, Gianna Lengyel, Guillaume Lample, Lucile Saulnier, L  lio Renard Lavaud, Marie-Anne Lachaux, Pierre Stock, Teven Le Scao, Thibaut Lavril, Thomas Wang, Timoth  e Lacroix, and William El Sayed. 2023. <a href="#">Mistral 7b</a> . <i>Preprint</i> , arXiv:2310.06825. | 942 |
| 886 |                                                                                                                                                                                                                                                                                                                                                                                                             | 943 |
| 887 |                                                                                                                                                                                                                                                                                                                                                                                                             | 944 |
| 888 |                                                                                                                                                                                                                                                                                                                                                                                                             | 945 |
| 889 |                                                                                                                                                                                                                                                                                                                                                                                                             | 946 |
| 890 |                                                                                                                                                                                                                                                                                                                                                                                                             | 947 |
| 891 |                                                                                                                                                                                                                                                                                                                                                                                                             | 948 |
| 892 |                                                                                                                                                                                                                                                                                                                                                                                                             | 949 |
| 893 | Cheng Jiayang, Chunkit Chan, Qianqian Zhuang, Lin Qiu, Tianhang Zhang, Tengxiao Liu, Yangqiu Song, Yue Zhang, Pengfei Liu, and Zheng Zhang. 2024. ECON: On the detection and resolution of evidence conflicts. In <i>Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing</i> , pages 7816–7844, Miami, Florida, USA. Association for Computational Linguistics.          | 950 |
| 894 |                                                                                                                                                                                                                                                                                                                                                                                                             | 951 |
| 895 |                                                                                                                                                                                                                                                                                                                                                                                                             | 952 |
| 896 |                                                                                                                                                                                                                                                                                                                                                                                                             | 953 |
| 897 |                                                                                                                                                                                                                                                                                                                                                                                                             | 954 |
| 898 |                                                                                                                                                                                                                                                                                                                                                                                                             | 955 |
| 899 |                                                                                                                                                                                                                                                                                                                                                                                                             | 956 |
| 900 |                                                                                                                                                                                                                                                                                                                                                                                                             | 957 |
| 901 | Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich K  ttler, Mike Lewis, Wen-tau Yih, Tim Rock-t  schel, Sebastian Riedel, and Douwe Kiela. 2020. Retrieval-augmented generation for knowledge-intensive NLP tasks. In <i>Advances in Neural Information Processing Systems</i> .                                                                      | 958 |
| 902 |                                                                                                                                                                                                                                                                                                                                                                                                             | 959 |
| 903 |                                                                                                                                                                                                                                                                                                                                                                                                             | 960 |
| 904 |                                                                                                                                                                                                                                                                                                                                                                                                             | 961 |
| 905 |                                                                                                                                                                                                                                                                                                                                                                                                             | 962 |
| 906 |                                                                                                                                                                                                                                                                                                                                                                                                             | 963 |
| 907 |                                                                                                                                                                                                                                                                                                                                                                                                             | 964 |
| 908 | Siyi Liu, Qiang Ning, Kishalay Halder, Zheng Qi, Wei Xiao, Phu Mon Htut, Yi Zhang, Neha Anna John, Bonan Min, Yassine Benajiba, and Dan Roth. 2025. Open domain question answering with conflicting contexts. In <i>Findings of the Association for Computational Linguistics: NAACL 2025</i> , pages 1838–1854, Albuquerque, New Mexico. Association for Computational Linguistics.                        | 965 |
| 909 |                                                                                                                                                                                                                                                                                                                                                                                                             | 966 |
| 910 |                                                                                                                                                                                                                                                                                                                                                                                                             | 967 |
| 911 |                                                                                                                                                                                                                                                                                                                                                                                                             | 968 |
| 912 |                                                                                                                                                                                                                                                                                                                                                                                                             | 969 |
| 913 |                                                                                                                                                                                                                                                                                                                                                                                                             | 970 |
| 914 |                                                                                                                                                                                                                                                                                                                                                                                                             | 971 |
| 915 |                                                                                                                                                                                                                                                                                                                                                                                                             | 972 |
| 916 | Yu Liu, Wenxiao Zhang, Diandian Guo, Cong Cao, Fangfang Yuan, Qiang Sun, Yanbing Liu, Jin B. Hong, and Zhiyuan Ma. 2026. <a href="#">CRAFT: Calibrated reasoning with answer-faithful traces via reinforcement learning for multi-hop question answering</a> . <i>Preprint</i> , arXiv:2602.01348.                                                                                                          | 973 |
| 917 |                                                                                                                                                                                                                                                                                                                                                                                                             | 974 |
| 918 |                                                                                                                                                                                                                                                                                                                                                                                                             |     |
|     | Llama Team, AI at Meta. 2024. <a href="#">The llama 3 herd of models</a> . <i>Preprint</i> , arXiv:2407.21783.                                                                                                                                                                                                                                                                                              |     |
|     | Yifei Ming, Senthil Purushwalkam, Shrey Pandit, Zixuan Ke, Xuan-Phi Nguyen, Caiming Xiong, and Shafiq Joty. 2025. FaithEval: Can your language model stay faithful to context, even if “the moon is made of marshmallows”. In <i>The Thirteenth International Conference on Learning Representations</i> .                                                                                                  |     |
|     | Aashiq Muhamed, Leonardo F. R. Ribeiro, Markus Dreyer, Virginia Smith, and Mona T. Diab. 2025. <a href="#">RefusalBench: Generative evaluation of selective refusal in grounded language models</a> . <i>Preprint</i> , arXiv:2510.10390.                                                                                                                                                                   |     |
|     | Ori Ram, Yoav Levine, Itay Dalmedigos, Dor Muhlgay, Amnon Shashua, Kevin Leyton-Brown, and Yoav Shoham. 2023. In-context retrieval-augmented language models. <i>Transactions of the Association for Computational Linguistics</i> , 11:1316–1331.                                                                                                                                                          |     |
|     | Freda Shi, Xinyun Chen, Kanishka Misra, Nathan Scales, David Dohan, Ed Chi, Nathanael Sch  rli, and Denny Zhou. 2023. <a href="#">Large language models can be easily distracted by irrelevant context</a> . <i>Preprint</i> , arXiv:2302.00093.                                                                                                                                                            |     |
|     | Maojia Song, Shang Hong Sim, Rishabh Bhardwaj, Hai Leong Chieu, Navonil Majumder, and Soujanya Poria. 2025. Measuring and enhancing trustworthiness of LLMs in RAG through grounded attributions and learning to refuse. In <i>The Thirteenth International Conference on Learning Representations</i> . ArXiv:2409.11242.                                                                                  |     |
|     | Zhaochen Su, Jun Zhang, Xiaoye Qu, Tong Zhu, Yanshu Li, Jiashuo Sun, Juntao Li, Min Zhang, and Yu Cheng. 2024. <a href="#">Conflictbank: A benchmark for evaluating the influence of knowledge conflicts in LLMs</a> . <i>Preprint</i> , arXiv:2408.12076.                                                                                                                                                  |     |
|     | Tu Vu, Mohit Iyyer, Xuezhi Wang, Noah Constant, Jerry Wei, Jason Wei, Chris Tar, Yun-Hsuan Sung, Denny Zhou, Quoc Le, and Thang Luong. 2024. Fresh-LLMs: Refreshing large language models with search engine augmentation. In <i>Findings of the Association for Computational Linguistics: ACL 2024</i> , pages 13697–13720, Bangkok, Thailand. Association for Computational Linguistics.                 |     |
|     | Alexander Wan, Eric Wallace, and Dan Klein. 2024. What evidence do language models find convincing? In <i>Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)</i> , pages 7468–7484, Bangkok, Thailand. Association for Computational Linguistics.                                                                                              |     |
|     | Chenyu Wang, Yueyuan Li, Yingmin Liu, and Yang Shu. 2026. <a href="#">ConflictRAG: Detecting and resolving knowledge conflicts in retrieval augmented generation</a> . <i>Preprint</i> , arXiv:2605.17301.                                                                                                                                                                                                  |     |



Han Wang, Archiki Prasad, Elias Stengel-Eskin, and Mohit Bansal. 2025. [Retrieval-augmented generation with conflicting evidence](#). *Preprint*, arXiv:2504.13079.

Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc V. Le, and Denny Zhou. 2022. Chain-of-thought prompting elicits reasoning in large language models. In *Advances in Neural Information Processing Systems*.

H. Wu, C. Wang, C. Sun, C. Lu, Z. Zhang, and Y. Chen. 2026. [Conflict-aware RAG: Multi-stage learning with conflict signals for robust retrieval-augmented generation](#). In *Proceedings of the ACM Web Conference 2026*, pages 2114–2125.

Rongwu Xu, Zehan Qi, Zhijiang Guo, Cunxiang Wang, Hongru Wang, Yue Zhang, and Wei Xu. 2024a. Knowledge conflicts for LLMs: A survey. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 8541–8565, Miami, Florida, USA. Association for Computational Linguistics.

Yilong Xu, Jinhua Gao, Xiaoming Yu, Baolong Bi, Huawei Shen, and Xueqi Cheng. 2024b. [ALiCE: Evaluating positional fine-grained citation generation](#). *Preprint*, arXiv:2406.13375.

An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, et al. 2024. [Qwen2.5 technical report](#). *Preprint*, arXiv:2412.15115.

Zhilin Yang, Peng Qi, Saizheng Zhang, Yoshua Bengio, William W. Cohen, Ruslan Salakhutdinov, and Christopher D. Manning. 2018. HotpotQA: A dataset for diverse, explainable multi-hop question answering. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 2369–2380, Brussels, Belgium. Association for Computational Linguistics.

Hua Ye, Siyuan Chen, Ziqi Zhong, Canran Xiao, Hao-liang Zhang, Yuhua Wu, and Fei Shen. 2026. [Seeing through the conflict: Transparent knowledge conflict handling in retrieval-augmented generation](#). *Preprint*, arXiv:2601.06842.

Ori Yoran, Tomer Wolfson, Ori Ram, and Jonathan Berant. 2024. Making retrieval-augmented language models robust to irrelevant context. In *The Twelfth International Conference on Learning Representations*.

Wenhao Yu, Hongming Zhang, Xiaoman Pan, Peixin Cao, Kaixin Ma, Jian Li, Hongwei Wang, and Dong Yu. 2024. Chain-of-note: Enhancing robustness in retrieval-augmented language models. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 14672–14685, Miami, Florida, USA. Association for Computational Linguistics.

Michael J. Q. Zhang and Eunsol Choi. 2021. SituatedQA: Incorporating extra-linguistic contexts into

QA. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 7371–7387, Online and Punta Cana, Dominican Republic. Association for Computational Linguistics.

Qinggang Zhang, Zhishang Xiang, Yilin Xiao, Le Wang, Junhui Li, Xinrun Wang, and Jinsong Su. 2025. FaithfulRAG: Fact-level conflict modeling for context-faithful retrieval-augmented generation. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Vienna, Austria. Association for Computational Linguistics.

Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. 2023. Judging LLM-as-a-judge with MT-bench and chatbot arena. In *Advances in Neural Information Processing Systems 36: Datasets and Benchmarks Track*.

## A Resource Construction

### A.1 Benchmark composition by source and track

| Source           | Track     | <i>n</i>   | NC         | CI         | CO         | OI         | MI        |
|------------------|-----------|------------|------------|------------|------------|------------|-----------|
| QACC             | web-retr. | 145        | 69         | 46         | 14         | 7          | 9         |
| SituatedQA       | web-retr. | 141        | 43         | 50         | 6          | 32         | 10        |
| ConflictingQA    | web-retr. | 137        | 8          | 53         | 73         | 1          | 2         |
| TRUST-ALIGN      | refusal   | 128        | 57         | 45         | 13         | 13         | 0         |
| FreshQA          | web-retr. | 85         | 33         | 8          | 2          | 27         | 15        |
| WikiRevision     | curated   | 78         | 0          | 0          | 0          | 78         | 0         |
| HotpotQA         | curated   | 19         | 0          | 19         | 0          | 0          | 0         |
| HealthContradict | curated   | 2          | 1          | 0          | 1          | 0          | 0         |
| Misinformation   | curated   | 1          | 0          | 0          | 0          | 0          | 1         |
| <b>Total</b>     |           | <b>736</b> | <b>211</b> | <b>221</b> | <b>109</b> | <b>158</b> | <b>37</b> |

Table 5: **Benchmark composition by source and construction track.** NC = No conflict, CI = Complementary information, CO = Conflicting opinions/research outcomes, OI = Conflict due to outdated information, MI = Conflict due to misinformation. Track A is web-retrieved, Track B is the held-out refusal pool, Track C is curated conflict. The curated track is class-targeted: all 78 WikiRevision items are temporal-supersession cases, supplying 49% of that class, and all 19 HotpotQA items are complementary-evidence cases.

Table 5 gives the full breakdown referenced in §4. The three tracks are structurally distinguishable, which is why we report them separately rather than as one pool. Track A records carry live retrieval URLs across 1,431 distinct domains and five documents each. Track B records use placeholder URLs inherited from the upstream refusal resource and are uniformly five documents. Track C



records are internally homogeneous by source: all 78 Wikipedia-revision items have exactly four documents drawn from a single domain with revision-dated URLs and all 19 multi-hop items have exactly ten documents with no URLs.

## A.2 Retrieval funnel

The non-refusal track proceeds through the following stages (Table 6). Counts are as observed in the retained artifacts.

Retrieval requests twenty results per query and retains at most ten, with a query-focused window selected by a TAS-B dense bi-encoder (Hofstätter et al., 2021). The released records carry a deterministic five-document view of those ten, sampled by a seeded rule that takes two from the top five and three from the bottom five; this is why Table 1 reports a mean of 5.03 documents per benchmark record rather than ten. The reduction keeps the evidence set answerable while preventing the conflicting documents from always occupying the highest-ranked positions, which would otherwise let a model succeed by attending to rank alone.

| Stage                                                | Records       |
|------------------------------------------------------|---------------|
| Raw candidate rows from four source pools            | 20,330        |
| — normalised overlaps with the training question set | — 598         |
| — duplicate normalised queries                       | — 8,891       |
| Unique usable candidates                             | 10,841        |
| Source-balanced selection (seeded)                   | 2,000         |
| Retained with exactly ten documents                  | 1,878 (93.9%) |
| underfilled, excluded                                | 122           |
| Human first-pass reviewed (canonical cleaned corpus) | 1,221         |
| Selected non-refusal set                             | 800           |
| Consolidated preselection/review population          | 1,454         |
| Final answerable benchmark records                   | 608           |

Table 6: Non-refusal construction funnel, counts as observed in the retained artifacts.

Two discontinuities in this funnel have no retained per-record selector and we report them as observed rather than reconstructing a rule. First, 1,878 candidates were assigned for first-pass review but the canonical cleaned corpus contains 1,221 unique reviewed records. The remaining discontinuity is that the quality selection from the 1,454-record consolidated population down to the released benchmark is described by a documented predicate, acceptance, reviewer confidence, retrieval quality, evidence sufficiency, conflict clarity, query specificity, source reliability, relevant-document count and gold-answer feasibility, but no per-record script was retained. Relatedly, 105 of the 608 answerable benchmark records appear in the consolidated review population only as

consensus-completed entries rather than as first-pass reviews. The refusal track reduces from a 200-record curated subset to the 128 released records without a documented reduction rule, though a 128-record refusal-quality review stratum with complete agreement is retained.

## A.3 Verification of the released files

All statistics below are computed directly from the released JSONL rather than from construction logs (Table 7; Figures 3 and 4).

**An independent validity check on the refusal track.** Refusal-required examples are not merely *labelled* insufficient, the document-level adjudication, which was performed independently of the abstention label, confirms it. Decisive *supports* verdicts occur on 3.6% of refusal-track documents against 39.3% elsewhere, an eleven-fold difference.




Figure 3: Resource composition. Left: the three benchmark construction tracks. Centre: the curated track is class-targeted, supplying 49% of the temporal-supersession class. Right: the benchmark is deliberately harder than the training distribution, with fewer no-conflict and more outdated/misinformation cases.

| Property                                           | Train       | Benchmark   |
|----------------------------------------------------|-------------|-------------|
| Records                                            | 943         | 736         |
| Duplicate identifiers                              | 0           | 0           |
| Identifier overlap (train/val $\cap$ benchmark)    | 0           | 0           |
| Document/note identifier misalignments             | 0           | 0           |
| Canonical d1 . . dN ordering                       | 943/943     | 736/736     |
| Key fact on every non-irrelevant note              | 5,886/5,886 | 3,393/3,393 |
| Irrelevant notes carrying a key fact               | 0           | 0           |
| Notes retaining an annotator vote tally            | 6,642/6,642 | 3,701/3,701 |
| Unique source domains                              | 2,633       | 1,431       |
| Mean snippet length (words)                        | 231         | 222         |
| Quote is an exact contiguous span                  | 82.6%       | 81.1%       |
| Quote is snippet-anchored (exact/ellipsis/partial) | 95.6%       | 95.9%       |

Table 7: Verification of the released files, computed directly from the released JSONL rather than from construction logs.




Figure 4: Quality audit computed directly from the released files.

## B Annotation Pipeline

Annotation is stagewise. At each stage every committee member independently receives the same structured input and returns a structured output; the committee votes only on the central categorical decision for that stage and then adopts the associated natural-language explanation *intact* from the highest-weight member that supported the winning decision. The method therefore combines committee agreement on decisions with a single internally coherent reasoning bundle and never averages or

splices prose from different annotators.

Formally, for a decision field with candidate values  $v$ , each successfully represented annotator  $m$  contributes its configured weight  $w_m$  and the system selects

$$v^* = \arg \max_v \sum_m w_m \mathbf{1}[y_m = v],$$

with exact ties broken deterministically by a fixed candidate ordering, so that re-running the pipeline on identical inputs reproduces identical labels. Writing  $\mathcal{M}^* = \{m : y_m = v^*\}$  for the annotators on the winning side, the explanation bundle is adopted *intact* from  $m^* = \arg \max_{m \in \mathcal{M}^*} w_m$  rather than assembled field-by-field across annotators, so a verdict, its quote and its rationale always originate from one coherent interpretation. Every note retains `_vote_tally`, `_winner_model` and `_all_verdicts`, so any label can be traced to its vote.

Stage-1 output is constrained: a verdict from the closed set, a key fact anchored to a quoted span of at most 50 words, a verdict rationale and a URL-derived source-quality label. The benchmark-mode Stage-1 prompt differs from the ordinary prompt in one consequential way: older, contradictory or otherwise conflict-bearing evidence about the target is retained as *partially supports* rather than discarded as *irrelevant*, so that Stage 2 can see the evidence needed to diagnose temporal and scientific conflicts at all. Stage-2 votes the conflict label and the answerability flag; Stage-3 votes the abstention decision independently of Stage-2 answerability, which is why a complete analysis must inspect both fields.

Different retained artifacts were produced by different committee configurations and we report the configuration alongside the artifact rather than

applying a single current default retrospectively. Structural validators check schema conformance, label inventories and document/note alignment; they do not establish quote entailment, factual correctness or temporal ordering, which is why the human review of Appendix C is reported separately rather than as a property of the automated pipeline.

## C Human Review of the Resource

Two workflows cover the two populations and must not be conflated. Both are released with their raw per-reviewer files, assignment manifests, consolidated pair records and machine-readable agreement metrics.

### C.1 Training conflict-type review

All 658 base records were assigned to exactly two of three reviewers under a seeded, pair-balanced, label-interleaved allocation, giving 1,316 assignments with loads of 439/439/438 and reviewer-pair counts of 220/219/219. Reviewers saw the query, the retrieved snippets and the committee-assigned label, and could retain it or replace it, recording confidence and a reason. This is a label-validation protocol, not a keep/drop filter: no item is removed from the pool by review.

| Review field            | agree | raw    | $\kappa$ |
|-------------------------|-------|--------|----------|
| Reviewed conflict type  | 787   | 83.46% | 0.7694   |
| Retain-or-change action | 791   | 83.88% | 0.1341   |
| Changed-label (boolean) | 791   | 83.88% | 0.1341   |
| Reviewer confidence     | 607   | 64.37% | 0.0159   |

Table 8: **Training conflict-type review agreement** over the 943-record consolidated population. The five-way taxonomy field is the primary result. The low chance-corrected values on the remaining rows reflect strongly imbalanced marginals (most reviewers accepted the committee label and reported high confidence), not label-quality failures: raw agreement and  $\kappa$  answer different questions when one category dominates.

The primary result in Table 8 is the five-way conflict-type field: 787 of 943 pairs agree, giving 83.46% raw agreement and  $\kappa = 0.7694$ .

**What the review layer actually changed.** Relative to the original committee label, the first review retains it for 801/943 records (84.94%), the second for 895/943 (94.91%), and both reviews retain it for 772/943 (81.87%). The review layer is therefore doing real work: roughly one record in five is contested by at least one reviewer.

**Reviewer labels and released labels are different quantities.** The consolidation deliberately does not overwrite the released `conflict_type`. The two distributions consequently differ, and we report both rather than implying the review ratified the release. Reviewers’ first-pass labels give No conflict 373, Complementary 289, Outdated 130, Conflicting opinions 120 and Misinformation 31, against released counts of 343, 286, 136, 158 and 20. The largest divergences are the conflicting-opinions and misinformation classes. Any analysis that treats the released label as human-ratified would therefore overstate the evidence; we treat the review as a reliability estimate on the taxonomy, not as a corrected gold standard.

### C.2 Benchmark preselection review

Seven reviewers were allocated the 1,878 exact-ten candidate set under seeded source balancing, with non-uniform queues of 150/150/300/300/326/326/326 and a source mix of ConflictingQA 258, FreshQA 394, SituatedQA-temporal 413, QACC 407 and SituatedQA-geographic 406. Each reviewer recorded an acceptance decision, a preliminary conflict type, confidence, retrieval quality, evidence sufficiency, conflict clarity, query specificity, source reliability, a relevant-document count bin, gold-answer feasibility and free-text notes.

| Review field              | agree | raw    | $\kappa$ |
|---------------------------|-------|--------|----------|
| Preliminary conflict type | 1378  | 94.77% | 0.9217   |
| Retention decision        | 1428  | 98.21% | 0.9228   |
| Retrieval quality         | 1436  | 98.76% | 0.9550   |
| Evidence sufficiency      | 1436  | 98.76% | 0.9600   |
| Conflict clarity          | 1441  | 99.11% | 0.9660   |
| Query specificity         | 1449  | 99.66% | 0.9875   |
| Source reliability        | 1446  | 99.45% | 0.9881   |
| Relevant-document count   | 1426  | 98.07% | 0.9527   |
| Gold-answer feasibility   | 1440  | 99.04% | 0.9669   |
| Reviewer confidence       | 1445  | 99.38% | 0.9767   |

Table 9: **Benchmark preselection agreement**, all ten recorded review fields over the 1,454-record consolidated population. Every field exceeds  $\kappa = 0.92$ . The primary result is the five-way conflict-type field. Because the second pass is sequential (Appendix C), these figures measure verification consistency rather than blind double annotation.

Every one of the ten recorded fields in Table 9 exceeds  $\kappa = 0.92$ , with the five-way conflict-type field at 94.77% raw agreement and  $\kappa = 0.9217$  (1,378 of 1,454 pairs agreeing).



**The second pass is sequential.** The second reviewer sees the complete first review and may accept, edit or reject it; across the 800 selected records there are 673 accept-as-is and 127 edited actions. These figures therefore measure verification consistency and must not be presented as blind independent double annotation. The second pass was nonetheless able to dissent: it introduced label values the first pass never used, including 13 *Unsure* conflict-type judgments and 15 reject or borderline-reject retention decisions.

**Reconstruction caveat.** Three second-review output files were lost and were reconstructed after the corresponding reviewers confirmed that every first review in their queue had been accepted as correct (114, 115 and 114 records). Those rows are therefore attested rather than independently recorded, and the reconstruction manifest is released alongside the review data.

### C.3 Refusal-quality validation

The 128 released refusal items were reviewed as a separate stratum against four dedicated fields in addition to the ten preselection fields. Agreement is *unanimous on all fourteen*: both reviewers judged all 128 items to require refusal, to have a valid abstention ground truth, to carry a high-quality evidence-gap rationale, and to be a valid refusal benchmark item; they also agreed unanimously that evidence is insufficient and that no defensible gold answer is possible for any of them. Because each of these fields is single-valued across the stratum,  $\kappa$  is undefined and we report exact agreement rather than substituting a coefficient.

This complements the automated check in Appendix A: decisive *supports* verdicts occur on 3.6% of refusal-track documents against 39.3% elsewhere. The refusal targets are thus corroborated twice, once by document-level adjudication performed independently of the abstention label and once by unanimous human review.

## D Fine-Tuning: Recipe Chronology

### D.1 Optimisation and checkpoint selection

All models are fine-tuned with QLoRA (Dettmers et al., 2023; Hu et al., 2021) at rank 32,  $\alpha = 64$  and dropout 0.05, with learning rate  $2 \times 10^{-4}$ , effective batch 16, maximum sequence length 12,288 and NEFTune  $\alpha = 5.0$  (Jain et al., 2023). Training minimises the standard token-level cross-entropy of the serialised trace and final response given the

prompt and retrieved documents; no auxiliary loss term is added. Checkpoints are selected on a held-out validation split by a composite score that sums document-verdict accuracy, response-format validity and abstention correctness, penalising false abstentions so that a model cannot score well by refusing answerable questions. A *recipe* is one fixed training mixture: the set and proportion of task views and prompt variants a run is trained on. Qwen2.5-7B and Qwen2.5-32B use recipe K; Llama-3.1-8B and Mistral-7B use recipe L.

The recipe was developed by error-directed revision rather than search. Table 10 names each intervention and its measured outcome, including the two steps that did not work.

Here “n/r” marks a step whose exact mixture size was not retained. Row counts are repeated supervised *views*, not independent questions: recipe L draws on 910 source examples. Three cautions apply to any causal reading. J, K and L each change data coverage, drill composition and sample weighting together, so they are principled recipe revisions rather than single-factor ablations. The source-hygiene intervention explored at step G is *not* present in the K/L mixtures and should not be attributed to the reported models. And the reported matrices use recipe K for the Qwen family and recipe L for Llama and Mistral, because no completed Qwen matrix exists for recipe L.

## E Inference Protocol

Generation is deterministic at temperature 0. Continuation budgets are profile-specific, strict 1,400 tokens base with a 3,200 cap, runtime 1,200/2,200, minimal 900/1,800, because a strict trace must enumerate every retrieved document and a uniform cap would selectively truncate long-context examples, confounding reasoning quality with output-space availability. Generation halts on the sentinel.

A retry re-runs the *identical* deterministic prompt with a larger token budget when the structural contract is unmet. It never resamples, appends corrective feedback, alters the retrieved documents or substitutes a prompt profile. The minimal profile has no structural retry gate by default, specifically to prevent the generation wrapper from turning a sparse-prompt condition into a hidden strict condition.

Every run writes both a raw and a sanitised output file. Sanitisation may canonicalise document order and identifiers, normalise label aliases,

| Step | Rows          | Intervention and outcome                                                                                                                                                                           |
|------|---------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| D    | 6,699         | Strict/runtime/minimal multi-task mixture. <i>Established that a trace survives minimal prompting.</i>                                                                                             |
| E    | n/r           | Broad conflict-label pressure. <b>Negative:</b> improved 32B conflict accuracy, degraded 7B. Capacity-sensitive.                                                                                   |
| F    | 7,308         | One targeted taxonomy-boundary drill per source example. Best historical conflict result at both scales; narrower and more effective than E.                                                       |
| G    | 8,526 / 8,715 | Model-specific probes: extra Stage-1 supervision (32B) and source-hygiene plus minimal-prompt pressure (7B). <b>Negative:</b> each recovered its target metric but cost conflict/contract quality. |
| J    | 10,674        | <i>Turning point.</i> Rebuilt train/validation geometry to include benchmark-like short answerable cases. Broad over-abstention substantially reduced; residual errors became localised.           |
| K    | 12,659        | +27 derived five-document answerable examples, document-boundary and partial-synthesis drills. Clean gain for 32B; genuine trade-off for 7B.                                                       |
| L    | 13,349        | +21 short answerable <i>no-conflict</i> examples, correcting a shortcut K risked introducing. Strong for Llama; Mistral still over-abstains.                                                       |

Table 10: **Fine-tuning recipe chronology.** Row counts are repeated supervised views, not independent questions. “n/r” marks a step whose exact mixture size was not retained.

clear key facts on irrelevant documents, drop out-of-range citations and append a missing sentinel where a usable trace exists. It *cannot* invent a missing trace, construct absent document judgments, write a conflict analysis or supply a missing answer; a record without a parseable trace passes through unchanged.

## F CATS Specification

### F.1 Component definitions

**Grounded Refusal (GR).** Whether the system answered exactly when the evidence permits answering. It is computed *deterministically* from the benchmark answerability label, never by a judge. We report refusal F1, where a true positive is the model declining to answer when the retrieved evidence does not support one.

**Behaviour Adherence (BA).** Whether the answer followed the response policy its conflict type requires, judged against a per-type rubric: a no-conflict item expects a direct answer, complementary evidence expects synthesis across documents, conflicting opinions expect both positions represented without adjudication, outdated information expects the current value with the superseded one flagged, and misinformation expects explicit correction. The rubric is reproduced verbatim in Appendix O.

**Factual Grounding (FG).** Whether each extracted claim is both supported by an eligible retrieved document *and* cited. Only documents whose *gold* per-document verdict is *supports* or *partially supports* are eligible, so an answer that is true in the world but unsupported by the retrieved

set scores zero by construction. Support may be distributed across a document pair.

**Single-Truth Recall (STR).** Whether the response explicitly adopts the target answer as its own conclusion, rather than merely quoting it, presenting it as one of several possibilities, or attributing it to a source without endorsing it. It is evaluated only for items where a single correct conclusion exists: no conflict, complementary information, outdated information and misinformation ( $n = 487$ ). Conflicting-opinion items are excluded because no such conclusion exists for them.

### F.2 Applicability and completeness

For example  $i$ , let  $A_i \in \{0, 1\}$  be gold answerability,  $\hat{A}_i$  the parsed answer/refuse decision and  $t_i \in \{1, \dots, 5\}$  the conflict type. For a metric  $M$ , write  $\mathcal{I}_M$  for its applicable examples and  $n_M = |\mathcal{I}_M|$ . Every dataset-level component below is an example-macro average over that set, never an average over placeholder zeros. Correct refusals carry no answer-content metrics, so

$$\mathcal{I}_{BA} = \mathcal{I}_{FG} = \{i : \neg(A_i = 0 \wedge \hat{A}_i = 0)\},$$

$$\mathcal{I}_{AQ} = \mathcal{I}_{FG},$$

and  $\mathcal{I}_{STR} \subseteq \mathcal{I}_{FG}$  additionally requires a gold answer and a single-truth conflict type. An *incorrect* refusal remains in these denominators with zero extracted content; it is not silently excluded. A headline aggregate is published only when every applicable example has a score and every represented type has a value; otherwise the run is marked incomplete rather than silently changing the denominator.

| Conflict type                             | Definition                                                                             | Expected behaviour                                                                                                                          |
|-------------------------------------------|----------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------|
| No conflict                               | All sources provide consistent and aligned information.                                | Return a unified, concise answer that synthesises the content accurately.                                                                   |
| Complementary information                 | Sources provide different but non-contradictory pieces that complete each other.       | Integrate all valid details into a richer, more comprehensive response.                                                                     |
| Conflicting opinions or research outcomes | Sources disagree due to different viewpoints, experimental results or interpretations. | Present each viewpoint clearly, describe the nature of the disagreement, and avoid taking a side unless the evidence strongly supports one. |
| Outdated information                      | Some sources contain older data or conclusions that may no longer be reliable.         | Prioritise newer, verified information while explicitly noting which sources are outdated.                                                  |
| Misinformation                            | A source provides factually incorrect, misleading or fabricated claims.                | Reject or correct the misinformation, give the verified facts, and briefly explain why the claim is wrong.                                  |

Table 11: **The five conflict types**, their definitions and the response behaviour each requires. Stage 2 assigns exactly one label per example; Behaviour Adherence (§6) scores the response against the expected behaviour for the assigned type.

### F.3 Grounded Refusal

The per-example decision gate is  $g_i = \mathbf{1}[\hat{A}_i = A_i]$ .  
With “answered” as the positive class,

$$TP = \sum_i \mathbf{1}[\hat{A}_i=1, A_i=1],$$

$$FP = \sum_i \mathbf{1}[\hat{A}_i=1, A_i=0],$$

$$FN = \sum_i \mathbf{1}[\hat{A}_i=0, A_i=1],$$

$$TN = \sum_i \mathbf{1}[\hat{A}_i=0, A_i=0],$$

giving  $P_{\text{ans}} = TP/(TP+FP)$ ,  $R_{\text{ans}} = TP/(TP+FN)$  and  $F_{1,\text{ans}} = 2P_{\text{ans}}R_{\text{ans}}/(P_{\text{ans}} + R_{\text{ans}})$ , with a zero denominator yielding zero. The refusal-positive orientation swaps  $(TP, FP, FN, TN) \rightarrow (TN, FN, FP, TP)$ :

$$P_{\text{ref}} = \frac{TN}{TN + FN}, \quad R_{\text{ref}} = \frac{TN}{TN + FP},$$

and analogously  $F_{1,\text{ref}}$ . Decision accuracy is  $(TP + TN)/N$ . Both orientations are reported because a model can have near-perfect answer recall while refusing far too rarely, and because the benchmark is 82.6% answerable, so answering everything already yields  $608/736 = 0.826$  accuracy (§6).

### F.4 Committee voting rule for BA and STR

The active local committee is Qwen3.5-397B-A17B, Mistral Small 4 and DeepSeek-R1-Distill-32B with priorities  $p_j = (6, 3, 2)$ . On a binary judgment, valid judge  $j$  emits vote  $z_{ij} \in \{0, 1\}$  and confidence  $c_{ij} \in [0, 1]$ , contributing effective weight and vote masses

$$a_{ij} = p_j \max(c_{ij}, 0.01),$$

$$W_i^+ = \sum_{j:z_{ij}=1} a_{ij}, \quad W_i^- = \sum_{j:z_{ij}=0} a_{ij}.$$

The decision and the continuous positive support are

$$d_i = \mathbf{1}[W_i^+ > W_i^-], \quad u_i = \frac{W_i^+}{W_i^+ + W_i^-}.$$

A weighted tie is negative. Timeout, API and parse failures are removed before forming these sums; an all-judge failure is recorded as an evaluation failure, not interpreted as a confident negative.

### F.5 Behaviour Adherence

Here  $z_{ij}$  asks whether the response follows the type-specific behaviour rubric (Appendix O). The reported component is binary,

$$h_i = d_i, \quad BA = \frac{1}{n_{BA}} \sum_{i \in \mathcal{I}_{BA}} h_i,$$

whereas the aggregate uses the continuous support  $b_i = u_i$ . The distinction preserves disagreement: a narrow weighted majority and a unanimous panel both give  $h_i = 1$  but different  $b_i$ .

### F.6 Factual Grounding

Let  $\mathcal{D}_i^+$  be the documents whose gold per-document verdict is *supports* or *partially supports*. The deterministic extractor returns at most eight claims; for claim  $k$ , let  $\mathcal{C}_{ik}$  be its deduplicated cited-document set. FG uses *raw priority* rather than confidence, since its judge schema emits no per-document confidence. With  $\mathcal{J}_{ik}$  the valid judges,

$$P_{ik} = \sum_{j \in \mathcal{J}_{ik}} \max(1, p_j), \quad \tau_{ik} = \frac{P_{ik}}{2},$$

1453

1454

$$m_{ik} = \begin{cases} 2, & |\mathcal{J}_{ik}| > 1, \\ 1, & |\mathcal{J}_{ik}| = 1. \end{cases}$$

1455

1456

1457

1458

For eligible document  $d$ , let  $V_{ikd} = \sum_j \mathbf{1}[j \text{ names } d]$  and  $P_{ikd} = \sum_j \max(1, p_j) \mathbf{1}[j \text{ names } d]$ . The accepted claim-specific support set is

1459

$$\mathcal{S}_{ik} = \{d \in \mathcal{D}_i^+ : P_{ikd} > \tau_{ik} \wedge V_{ikd} \geq m_{ik}\}.$$

1460

1461

1462

1463

1464

1465

1466

For cross-document support, let  $x_{ikj}$  be judge  $j$ 's binary vote that two documents jointly support the claim, with mass  $X_{ik} = \sum_j \max(1, p_j) x_{ikj}$  and count  $U_{ik} = \sum_j x_{ikj}$ ; the same thresholds give  $\chi_{ik} = \mathbf{1}[X_{ik} > \tau_{ik} \wedge U_{ik} \geq m_{ik}]$  and accepted combination set  $\mathcal{D}_{ik}^\times$ . A claim earns credit only when accepted support is *actually cited*:

1467

1468

1469

$$y_{ik} = \mathbf{1}[(\mathcal{C}_{ik} \cap \mathcal{S}_{ik} \neq \emptyset) \vee (\chi_{ik} = 1 \wedge \mathcal{C}_{ik} \cap \mathcal{D}_{ik}^\times \neq \emptyset)].$$

1470

For  $K_i$  extracted claims,

1471

$$f_i = \begin{cases} K_i^{-1} \sum_{k=1}^{K_i} y_{ik}, & K_i > 0, \\ 0, & K_i = 0, \end{cases}$$

1472

1473

$$\text{FG} = \frac{1}{n_{\text{FG}}} \sum_{i \in \mathcal{I}_{\text{FG}}} f_i.$$

1474

1475

1476

1477

1478

1479

1480

1481

1482

The outer example-macro average prevents a long response from gaining weight merely by yielding more claims. Under the full (6, 3, 2) panel  $\tau = 5.5$ : the priority-6 judge clears the mass threshold alone but fails the two-judge corroboration rule, while priorities 3 + 2 fail the mass threshold, so no single judge can unilaterally certify a document and the anchor judge must be joined by at least one other (see Limitations).

1483

## F.7 Self-Truth Recall

1484

1485

1486

1487

1488

1489

1490

1491

1492

STR applies when a gold answer exists and  $t_i \in \{1, 2, 4, 5\}$ ; type 3 is excluded because conflicting-opinion cases need not have one canonical proposition. Let  $\mathcal{G}_i$  contain  $M_i$  gold answer items. Applying the same weighted rule to each item  $\ell$ , let  $e_{i\ell}$  be the committee decision that the response asserts that item as its own conclusion. For a negative decision the positive side is the minority, so define a partial-credit indicator

1493

$$p_{i\ell} = \mathbf{1}\left[e_{i\ell} = 0 \wedge \frac{W_{i\ell}^+}{W_{i\ell}^+ + W_{i\ell}^-} \geq 0.30\right],$$

1494

1495

$$E_i = \sum_{\ell} e_{i\ell}, \quad P_i = \sum_{\ell} p_{i\ell}.$$

The per-example and dataset scores are

1496

$$r_i = \begin{cases} 0, & M_i = 0 \text{ or empty candidate,} \\ 1, & E_i > 0, \\ \min\left(1, \frac{1}{2} \frac{P_i}{M_i}\right), & E_i = 0, P_i > 0, \\ 0, & \text{otherwise,} \end{cases}$$

1497

1498

1499

$$\text{STR} = \frac{1}{n_{\text{STR}}} \sum_{i \in \mathcal{I}_{\text{STR}}} r_i.$$

Quoting the gold item, listing it as one possibility among several, or attributing it to a source without adopting it does not satisfy  $e_{i\ell} = 1$ .

1500

1501

1502

## F.8 Aggregation and its structural properties

1503

The choice of aggregate is consequential: Figure 5 shows that a flat mean hides the refusal-calibration effect that both CATS aggregates surface, and Figure 6 decomposes a headline score into the part earned on answerable content and the part earned on the refusal decision alone.

1504

1505

1506

1507

1508

1509

The three-step construction is as follows. Grounding and recall are averaged *geometrically* into an answer-quality term; that term is averaged *harmonically* with behaviour adherence; and the result is multiplied by the grounded-refusal decision, which is 1 when the answer/refuse choice was correct and 0 otherwise. Both means are used because they penalise imbalance: a high score on one component cannot compensate for a low score on another, and a wrong decision sets the example to zero regardless of everything else. CATS-Prevalence is the plain mean over the benchmark; CATS-Balanced weights each conflict type equally.

1510

1511

1512

1513

1514

1515

1516

1517

1518

1519

1520

1521

1522

Per example, grounding and recall first form the answer-quality pillar

1523

1524

$$q_i = \begin{cases} \sqrt{f_i r_i}, & i \in \mathcal{I}_{\text{FG}} \cap \mathcal{I}_{\text{STR}}, \\ f_i, & i \in \mathcal{I}_{\text{FG}} \setminus \mathcal{I}_{\text{STR}}, \\ \text{N/A}, & i \notin \mathcal{I}_{\text{FG}}, \end{cases}$$

1525

1526

1527

$$\text{AQ} = \frac{1}{n_{\text{AQ}}} \sum_{i \in \mathcal{I}_{\text{AQ}}} q_i.$$

Let  $H(x, y) = 2xy/(x + y)$  for  $x, y > 0$  and  $H(x, y) = 0$  otherwise. The complete per-example

1528

1529






Figure 5: The choice of aggregate determines the conclusion. A flat mean over the four components hides the refusal-calibration effect almost entirely for Qwen2.5-32B, while CATS-Prevalence and CATS-Balanced both surface it. This is why we report components rather than a single number.




Figure 6: Refusal credit inside the aggregate. Left: CATS-Prevalence split into the part earned on answerable examples and the part earned on correct refusals, which are scored on their decision alone. Right: plotting the answerable-only score against the headline places every trace-tuned point below the diagonal, showing the headline is inflated by decision credit.

score is

$$s_i = \begin{cases} g_i, & A_i = 0, \\ g_i H(b_i, q_i), & A_i = 1, b_i, q_i \text{ available}, \\ g_i b_i, & A_i = 1, b_i \text{ only}, \\ g_i q_i, & A_i = 1, q_i \text{ only}, \\ g_i, & A_i = 1, \text{ neither available}. \end{cases}$$

The last three are auditable compatibility branches; complete benchmark runs use only the first two. The decision gate means a wrong answer/refuse decision cannot be rescued downstream, and the harmonic fusion penalises an answer that is weak on either behaviour or answer quality rather than letting a strong one mask it.

To define the two headline summaries, let  $\mathcal{I}_t^A = \{i : t_i = t, A_i = 1\}$  and  $\mathcal{I}_t^R = \{i : t_i = t, A_i = 0\}$ , with subgroup means

$$T_t^A = \frac{1}{|\mathcal{I}_t^A|} \sum_{i \in \mathcal{I}_t^A} s_i, \quad T_t^R = \frac{1}{|\mathcal{I}_t^R|} \sum_{i \in \mathcal{I}_t^R} s_i.$$

The decision-balanced score for type  $t$  is

$$T_t^{\text{DB}} = \begin{cases} \frac{1}{2}(T_t^A + T_t^R), & |\mathcal{I}_t^A| > 0, |\mathcal{I}_t^R| > 0, \\ T_t^A, & |\mathcal{I}_t^A| > 0, |\mathcal{I}_t^R| = 0, \\ T_t^R, & |\mathcal{I}_t^A| = 0, |\mathcal{I}_t^R| > 0, \end{cases}$$

and the two reported aggregates are

$$\text{CATS-P} = \frac{1}{N} \sum_{i=1}^N s_i$$

$$\text{CATS-B} = \frac{1}{5} \sum_{t=1}^5 T_t^{\text{DB}}$$

CATS-P preserves benchmark prevalence. CATS-B gives each conflict type equal weight and, within each type where both decisions occur, gives the answerable and refusal-required subgroups equal weight. Both are secondary to the four components (§6).

We disclose three consequences.

**Refusal-required examples score decision correctness only.** Two correct refusals receive the same score regardless of how well either is justified. This is a scope choice, not an oversight: no consistently audited refusal-quality judgment exists across the runs.

**The type-balanced aggregate over-weights that decision.** Balancing answerable and refusal subgroups within each conflict type gives refusal-required examples 40% of the type-balanced score against 17.4% of the data. One type-3 refusal example carries roughly twelve times the influence of one type-1 answerable example and the misinformation class, 5% of the data, determines 20% of the score.

**There is a non-trivial floor.** A model with zero usable answer content but perfect refusal calibration scores exactly 0.400 on the type-balanced aggregate. We observed this empirically in the excluded rows of Appendix L, which is how we came to characterise it (Figure 7).

## G Human Evaluation Study

### G.1 Per-component reliability

Four reviewers independently judged 300 generated responses, each reviewed twice, spanning two model families, two training conditions, three prompt profiles and all five conflict types. These 300 responses form the agreement frame on which the human-human statistics are computed (Table 12).

The result is a per-component reliability ordering identical on two independent instruments,




Figure 7: Structural properties of the aggregate.

| Component | human-comm. |              | h.-h.    | judges       |
|-----------|-------------|--------------|----------|--------------|
|           | agr.        | $\kappa$     | $\kappa$ | $\alpha$     |
| FG        | 0.899       | <b>0.783</b> | 0.780    | <b>0.871</b> |
| STR       | 0.908       | <b>0.790</b> | 0.773    | 0.678        |
| BA        | 0.745       | 0.407        | 0.669    | 0.457        |

Table 12: **Evaluator validation against independent human review.** Rows are Factual Grounding, Single-Truth Recall and Behaviour Adherence. Columns are human-committee agreement and  $\kappa$ , human-human  $\kappa$ , and *judges*, the nominal  $\alpha$  among the three committee models. Human-committee uses all accepted reviewer rows; human-human and *judges* use the 300-response complete two-reviewer frame. The reliability ordering is identical on both instruments: grounding and recall are reliable, behaviour is not. Grounded Refusal is deterministic and therefore not judged.

human-committee agreement and committee-internal agreement, and it bounds what we may claim. Factual Grounding and Single-Truth Recall are reliable ( $\kappa = 0.78$  and  $0.79$  against humans; 89.9% and 90.8% raw agreement). Behaviour Adherence is not: an independent human panel agrees with the committee on only 74.5% of behaviour judgments ( $\kappa = 0.41$ ), and the three judges agree among themselves at  $\alpha = 0.46$ .

Consequently we treat behaviour claims differently from grounding claims. All twelve observed BA differences lie below the human-committee disagreement rate, so while the improvement is statistically consistent (11/12 cells,  $p = 0.006$ ), we state it qualitatively and decline to interpret individ-

ual BA margins, particularly in the misinformation regime.

The agreement frame is **300 generated responses**, each independently reviewed by two of four reviewers, drawn from a  $2 \times 3 \times 2$  design (two model families  $\times$  three prompt profiles  $\times$  untuned/fine-tuned), stratified by conflict type, with correct refusals excluded before assignment. Reviewers judged behaviour adherence against the per-type rubric, selected supporting documents for each displayed claim and judged whether the response asserted the gold target; metric pages were separated so that a factual impression could not silently change a behaviour label.

A further 50 responses received a single review, because one reviewer completed half an assigned queue; these carry no human-human comparison and are excluded from the agreement frame, but their judgments are retained where a single review suffices. No judgment is imputed. Reported statistics therefore name their denominator explicitly, and the denominators differ by construction: the *response-level* human-human frame is 300, whereas *judgment-level* human-committee comparison pools every accepted review and so is 650 ( $300 \times 2$  complete pairs plus 50 singles). Claim-level grounding is counted in claim units rather than responses (477 human-human, 1,029 human-committee), and the human-consensus frame is 256 behaviour and 194 strict-recall samples.

Coverage caveats. The study covers two of the four model families and the end-to-end condition only. Correct refusals were excluded before sampling, so the study says nothing about refusal wording quality. The human package extracts up to twelve claims per answer while the benchmark configuration uses eight, so grounding denominators are aligned in logic but not identical.

**Agreement is not uniform across the taxonomy.** Table 15 gives the per-conflict-type breakdown. Three patterns matter. Claim-level grounding is near-perfect on No-conflict items ( $\kappa = 0.972$  human-human) and weakest on Outdated ( $\kappa = 0.566$ ), where deciding whether a superseded document still supports a claim is genuinely ambiguous. Human behaviour agreement is lowest on Complementary ( $\kappa = 0.536$ ) and Outdated ( $\kappa = 0.505$ ). And the committee’s behaviour judgments track humans well on Conflicting opinions ( $\kappa = 0.641$ ) but poorly on Complementary ( $\kappa = 0.279$ ) and Misinformation ( $\kappa = 0.149$ ). Reported as a single

| Judge (behaviour)                                        | prio. | agr.           | $\kappa$     |
|----------------------------------------------------------|-------|----------------|--------------|
| anchor judge                                             | 6     | 0.743          | 0.405        |
| third judge                                              | 2     | 0.735          | 0.339        |
| second judge                                             | 3     | 0.643          | 0.229        |
| <b>full committee</b>                                    |       | <b>0.745</b>   | <b>0.407</b> |
| <i>Committee-internal agreement (300-response frame)</i> |       |                |              |
| Factual Grounding                                        |       | $\alpha=0.871$ | 0.871        |
| Single-Truth Recall                                      |       | $\alpha=0.678$ | 0.687        |
| Behaviour Adherence                                      |       | $\alpha=0.457$ | 0.470        |

Table 13: **Per-judge agreement with human reviewers, and committee-internal agreement.** Top: each judge’s behaviour agreement with human reviewers over 650 aligned judgment-level rows (the 300 two-reviewer responses plus 50 single-reviewed), alongside its configured priority. The anchor judge is empirically the best human-aligned of the three, and the full committee edges its best single member. Bottom: agreement among the three judges on the 300-response frame; the final column is mean pairwise  $\kappa$ .

number, behaviour reliability would hide the fact that it is concentrated in two regimes.

**Per-judge and committee-internal agreement.** Table 13 compares each judge separately against human reviewers over the 650 aligned judgment-level rows. The anchor judge agrees best (0.743 agreement,  $\kappa = 0.405$ ), the third judge next (0.735,  $\kappa = 0.339$ ) and the second judge least (0.643,  $\kappa = 0.229$ ); the full weighted committee reaches  $\kappa = 0.407$ , marginally above its best individual member. Committee-internal agreement on the 300-response frame is  $\alpha = 0.871$  for grounding, 0.678 for recall and 0.457 for behaviour, with the behaviour spread concentrated between judge pairs (pairwise  $\kappa$  ranges from 0.368 to 0.611).

**Consensus frame.** Restricting to the samples where both reviewers agree and comparing that consensus against the committee raises agreement as expected: behaviour 0.789 ( $\kappa = 0.483$ ,  $n = 256$ ), strict recall 0.948 ( $\kappa = 0.876$ ,  $n = 194$ ) and claim-level grounding 0.949 ( $\kappa = 0.888$ ,  $n = 428$ ). Consensus filtering improves interpretability but retains easier cases disproportionately, so we report it alongside rather than instead of the row-level frame. The misinformation behaviour stratum does not improve under consensus (0.588,  $\kappa = 0.165$ ), which confirms that its weakness is not merely reviewer noise.

**Alignment integrity.** Claim-level comparisons require pairing human and committee claim in-

ventories. All three analyses report *zero* claim-alignment issues, and all 650 accepted human reviews matched a committee judgment with no missing matches, so no unit was dropped or approximately aligned (Figure 8).




Figure 8: Left: paired effect of trace supervision across the twelve end-to-end cells, with the distribution-free noise floor shaded. Right: per-component reliability from the human study.

## H Citation Analysis

Few-shot Chain-of-Thought abstains at 32/128 for Llama-3.1-8B, 38/128 for Qwen2.5-7B and 85/128 for Mistral-7B, a 25–66% spread. The exception is worth naming rather than averaging away: Mistral-7B is also the most abstention-prone model in the baseline end-to-end condition, so its few-shot behaviour looks less like acquired judgment than like a prior toward refusing that the exemplars happen to license. Chain-of-Note has the poorest citation hygiene of any method examined, with a hard-negative citation rate roughly five times that of the trace-tuned models. Figure 11 contrasts the two prompting strategies with trace supervision on answer content and on refusal calibration.

We analyse cited-document behaviour with metrics that are not part of CATS and do not alter any CATS component or aggregate; they are supplementary post-hoc diagnostics computed from the same claim-level extraction used by Factual Grounding, partitioned into a recall side (does the

model cite the support that exists?) and a precision side (is what it cites the right evidence?).

Table 14 defines each metric; Table 16 reports the changes.

**Unit of analysis.** Every metric below is computed over *distinct cited document identifiers per claim*, pooled across claims, not raw citation mentions: if a claim cites the same document twice, that citation is counted once for that claim. Micro-averages pool these per-claim counts across every evaluated claim in the run rather than averaging a per-claim rate, so claims with more citations are not down-weighted relative to claims with fewer.

| Metric                                      | Formula                                                                              | Unit            |
|---------------------------------------------|--------------------------------------------------------------------------------------|-----------------|
| Claim-spec. citation prec. $\uparrow$       | $\frac{\text{approved}}{\text{all cited}}$                                           | cited doc       |
| Gold-positive precision $\uparrow$          | $\frac{\text{gold-pos.}}{\text{all cited}}$                                          | cited doc       |
| Citation sufficiency $\uparrow$             | $\frac{\text{claims w/} \geq 1 \text{ cite}}{\text{claims w/ support}}$              | supported claim |
| Soft-extra citation rate $\downarrow$       | $\frac{\text{gold-pos., not appr.}}{\text{all cited}}$                               | cited doc       |
| Hard-negative citation rate $\downarrow$    | $\frac{\text{not gold-pos.}}{\text{all cited}}$                                      | cited doc       |
| Strict grounded, gold-clean $\uparrow$      | $\frac{\text{gr.} \wedge \text{cit.} \wedge \text{no hard-neg.}}{\text{all claims}}$ | extras OK       |
| Strict grounded, committee-clean $\uparrow$ | $\frac{\text{gr.} \wedge \text{cit.} \wedge \text{clean}}{\text{all claims}}$        | nothing OK      |

Table 14: **Citation-metric definitions.** “Approved” means the cited document is in the committee’s claim-specific support set; “gold-positive” means its gold per-document verdict is *supports* or *partially supports*, independent of which specific claim is being checked; “clean” means no soft-extra and no hard-negative citation. The two strict-grounded rates share a base condition (committee-grounded, cited) and differ only in what they tolerate: the gold-clean rate forbids hard negatives but allows soft extras, the committee-clean rate forbids both, making it the single most demanding number in this section.  $\uparrow/\downarrow$  after the metric name marks whether higher or lower is the better direction. All fractions count distinct cited documents or evaluated claims as noted in the Unit column.

Only the recall side moves. The strict committee-clean metric, which requires a claim to be committee-supported, cited and free of both extraneous and incorrect citations, is positive in all twelve cells and is the single most demanding citation number we report.



| Conflict type        | human-human |          |             |          | human-committee |              |             |          |
|----------------------|-------------|----------|-------------|----------|-----------------|--------------|-------------|----------|
|                      | Behaviour   |          | FG (claims) |          | Behaviour       |              | FG (claims) |          |
|                      | agr.        | $\kappa$ | agr.        | $\kappa$ | agr.            | $\kappa$     | agr.        | $\kappa$ |
| No conflict          | 0.937       | 0.714    | 0.987       | 0.972    | 0.865           | <b>0.325</b> | 0.945       | 0.880    |
| Complementary        | 0.783       | 0.536    | 0.908       | 0.803    | 0.708           | <b>0.279</b> | 0.910       | 0.805    |
| Conflicting opinions | 0.862       | 0.725    | 0.887       | 0.749    | 0.820           | 0.641        | 0.888       | 0.759    |
| Outdated             | 0.817       | 0.505    | 0.800       | 0.566    | 0.746           | <b>0.338</b> | 0.825       | 0.606    |
| Misinformation       | 0.864       | 0.725    | 0.912       | 0.821    | 0.581           | <b>0.149</b> | 0.929       | 0.855    |
| <b>Overall</b>       | 0.853       | 0.669    | 0.897       | 0.780    | 0.745           | 0.407        | 0.899       | 0.783    |

Table 15: **Agreement by conflict type.** Bold marks  $\kappa < 0.40$ . Reliability is not uniform across the taxonomy: human reviewers agree least on Complementary and Outdated behaviour, and the committee’s behaviour judgments degrade sharply on Misinformation ( $\kappa = 0.149$ ) and Complementary ( $\kappa = 0.279$ ) while remaining strong on Conflicting opinions ( $\kappa = 0.641$ ). Claim-level grounding is reliable everywhere except Outdated. Conflicting opinions is excluded from Single-Truth Recall by construction and so does not appear in that metric’s strata.

| Metric                                | mean $\Delta$ | min    | max    |
|---------------------------------------|---------------|--------|--------|
| Citation sufficiency (support recall) | +0.363        | −0.008 | +0.869 |
| Strict grounded, gold-clean           | +0.358        | −0.002 | +0.811 |
| Strict grounded, committee-clean      | +0.339        | +0.119 | +0.651 |
| Gold-positive precision               | +0.003        | −0.004 | +0.014 |
| Soft-extra citation rate              | +0.006        | −0.058 | +0.074 |
| Hard-negative citation rate           | −0.003        | −0.014 | +0.004 |
| Claim-specific citation precision     | −0.003        | −0.077 | +0.055 |

Table 16: Citation-metric changes under trace supervision, fine-tuned minus baseline, across the twelve end-to-end cells.

Two further observations. Untuned models under minimal prompts post *higher* claim-specific citation precision than their own fine-tuned counterparts in all four cases (0.736–0.867 against 0.713–0.800) while citing almost nothing (recall 0.036–0.139): precision over a near-empty denominator, not a quality signal. And the recall-side metrics correlate with Factual Grounding at  $\rho = 0.93$ –0.98, so they are largely redundant with it, whereas every precision-side metric sits at  $|\rho| \leq 0.46$  against every CATS component and therefore carries information the framework does not otherwise capture. Behaviour Adherence is the CATS component least related to citation behaviour, which is an independent check that its orthogonality requirement holds in practice. Figure 9 plots the sufficiency and precision trajectories across the instruction ladder, and Figure 10 shows the precision paradox directly.

## I Bounding Factual Grounding

The *lower* bound includes the missed refusals; the *upper* bound restricts to the like-for-like population where the model answered *and* the gold label is answerable. The intervals separate in all four minimal-prompt cells and in only two others (Qwen2.5-7B runtime, Mistral-7B strict). Instruc-




Figure 9: Trace supervision fixes citation sufficiency, not citation precision.




Figure 10: The precision paradox: untuned precision is an artefact of abstention from citing. Every model moves down-and-right when fine-tuned.

tion volume is also not monotonically helpful for baseline models: grounding is higher under the shorter runtime prompt than the strict one in all four families, and in the extreme case Llama-3.1-8B reaches  $FG = 0.897$  while correctly refusing 2 of 128 refusal-required items, so its grounding is computed over a markedly easier population.

Reported Factual Grounding is averaged over  $fg\_n = 608 + \text{missed refusals}$ , so run types that refuse at different rates are not scored on the same population. Table 17 gives both bounds per cell: the reported value, which includes missed refusals, and the like-for-like interval restricted to examples the model answered and the gold label marks answerable. Figure 12 plots both, alongside the model-dependence of the denominator itself.

## J Additional Figures and Tables

### J.1 Behaviour Adherence by prompt condition

Figures 14 and 15 give the per-model Behaviour Adherence scores behind the summary in §8.1, separated by instruction strength. Two patterns are visible that the aggregate range hides. Under the strict prompt the gain is strongly ordered by how well the baseline already follows instructions: Mistral-7B starts lowest (0.660) and gains most (+0.133), Llama-3.1-8B gains +0.084, while the two Qwen baselines already sit at 0.789 and 0.811 and move +0.032 and −0.031. Only the first two clear the  $\pm 0.034$  noise floor, so on the strict prompt trace supervision is best read as closing a gap for weaker instruction-followers rather than as lifting every model.

| Model        | Prompt  | reported FG |       | like-for-like FG |             | sep. |
|--------------|---------|-------------|-------|------------------|-------------|------|
|              |         | base        | SFT   | base             | SFT         |      |
| Qwen2.5-7B   | strict  | 0.617       | 0.838 | [0.81,0.92]      | [0.86,0.87] | –    |
|              | runtime | 0.651       | 0.829 | [0.69,0.85]      | [0.85,0.86] | ✓    |
|              | minimal | 0.031       | 0.837 | [0.00,0.04]      | [0.86,0.88] | ✓    |
| Qwen2.5-32B  | strict  | 0.851       | 0.812 | [0.84,1.00]      | [0.85,0.85] | –    |
|              | runtime | 0.864       | 0.814 | [0.86,1.00]      | [0.85,0.85] | –    |
|              | minimal | 0.108       | 0.822 | [0.00,0.13]      | [0.86,0.86] | ✓    |
| Llama-3.1-8B | strict  | 0.705       | 0.852 | [0.75,0.94]      | [0.88,0.88] | –    |
|              | runtime | 0.897       | 0.849 | [0.88,1.00]      | [0.87,0.88] | –    |
|              | minimal | 0.074       | 0.857 | [0.00,0.09]      | [0.88,0.88] | ✓    |
| Mistral-7B   | strict  | 0.624       | 0.796 | [0.58,0.74]      | [0.87,0.87] | ✓    |
|              | runtime | 0.739       | 0.762 | [0.71,0.90]      | [0.83,0.83] | –    |
|              | minimal | 0.079       | 0.792 | [0.00,0.10]      | [0.86,0.86] | ✓    |

Table 17: **The Factual-Grounding denominator is model-dependent.**  $fg\_n = 608 + \text{missed refusals}$ , so a baseline model that rarely refuses is additionally scored on  $\sim 130$  refusal-required items whose evidence is insufficient by construction, while its false abstentions enter the denominator at  $FG = 0$ . The *like-for-like* columns bound FG on the population where the model answered *and* the gold label is answerable. A check mark marks cells where the fine-tuned interval lies strictly above the baseline interval: this holds in every minimal-prompt cell and in none of the strict cells.

| Model      | Min. FG |              | Spread |              |
|------------|---------|--------------|--------|--------------|
|            | base    | SFT          | base   | SFT          |
| Qwen-7B    | 0.031   | <b>0.837</b> | 0.442  | <b>0.026</b> |
| Qwen-32B   | 0.108   | <b>0.822</b> | 0.398  | <b>0.009</b> |
| Llama-8B   | 0.074   | <b>0.857</b> | 0.345  | <b>0.016</b> |
| Mistral-7B | 0.079   | <b>0.792</b> | 0.278  | <b>0.020</b> |

Table 18: **Internalization.** Factual Grounding under the minimal prompt, and the spread of CATS-Balanced across the strict/runtime/minimal ladder. Baseline grounding collapses without scaffolding and baseline behaviour swings an order of magnitude more across the ladder.

Under the minimal prompt the picture is more uniform. Every model gains (+0.031 to +0.062) and the baselines compress into a narrow band (0.727–0.747), because a forty-word instruction that never names the taxonomy removes the advantage the stronger instruction-followers had. The Qwen2.5-7B cell that regresses under strict prompting turns positive here. This is the same internalization effect §8.2 reports, seen on a single component: what the strict prompt supplies externally, trace supervision supplies internally, so the two conditions converge for tuned models and diverge for baselines. Because Behaviour Adherence is the least reliable CATS component ( $\kappa = 0.41$  against human review, §9), these per-cell margins should be read as directional rather than as point estimates.




Figure 11: External prompting comparators. Both prompting-only strategies reach competitive answer content but neither acquires refusal calibration: Chain-of-Note records 0/128 correct refusals on every model.




Figure 12: The Factual-Grounding denominator confound. Left: FG on the like-for-like population, with bars giving the bound and vertical marks the reported value; crosses mark cells where the intervals overlap and no gap is established. Right: the denominator  $fg\_n = 608 + \text{missed refusals}$  is itself model-dependent, so baseline models are scored on up to  $\sim 130$  extra items whose evidence is insufficient by construction.

The flatness is a property of the whole ladder, not only its endpoints: across all twenty model  $\times$  metric cells the intermediate runtime condition never falls more than 0.030 outside the strict–minimal interval for a trace-tuned model, inside the  $\pm 0.034$  noise floor. Single-truth recall shows the same pattern, with baseline spreads of 0.047–0.196. The result also reproduces on measurements never part of CATS: analysing cited-document behaviour separately (Appendix H), baseline citation sufficiency swings by 0.730 across the prompt ladder while fine-tuned sufficiency swings by 0.011, a factor of roughly seventy, on an instrument built after and independently of the framework whose claim it corroborates.

## K Oracle Ablations: Full Results

### K.1 Additivity, exceptions and baseline behaviour

Over the cells common to all three oracle conditions, the residual  $|\Delta_{\text{both}} - (\Delta_{\text{conflict}} + \Delta_{\text{notes}})|$  is at most 0.012 on every grounding metric. Behaviour is the single exception, with a residual of  $-0.014$ : supplying both oracles yields *less* behavioural gain than the conflict label alone, consistent with the long injected notes block competing for attention with the label the model is asked to reproduce.

Two caveats keep the claim honest. The effects are small in absolute terms, so it is their *consistency* across cells rather than their magnitude that carries the evidence. And supplying perfect Stage-1 and Stage-2 information still leaves most of the residual

| Metric                 | + gold conflict label |        |              | + gold per-doc notes |        |              | + both        |        |              |
|------------------------|-----------------------|--------|--------------|----------------------|--------|--------------|---------------|--------|--------------|
|                        | mean $\Delta$         | +/ $n$ | $p$          | mean $\Delta$        | +/ $n$ | $p$          | mean $\Delta$ | +/ $n$ | $p$          |
| Behaviour Adherence    | <b>+0.013</b>         | 8/9    | <b>0.039</b> | +0.008               | 10/12  | <b>0.039</b> | +0.011        | 9/12   | 0.146        |
| Factual Grounding      | -0.006                | 3/9    | 0.508        | <b>+0.047</b>        | 12/12  | <b>0.000</b> | <b>+0.046</b> | 12/12  | <b>0.000</b> |
| Answer Quality         | -0.009                | 2/9    | 0.180        | <b>+0.044</b>        | 12/12  | <b>0.000</b> | <b>+0.044</b> | 12/12  | <b>0.000</b> |
| Citation sufficiency   | -0.001                | 3/9    | 0.508        | <b>+0.028</b>        | 12/12  | <b>0.000</b> | <b>+0.030</b> | 12/12  | <b>0.000</b> |
| Strict grounded (cit.) | -0.007                | 3/9    | 0.508        | <b>+0.041</b>        | 12/12  | <b>0.000</b> | <b>+0.044</b> | 12/12  | <b>0.000</b> |
| Citation precision     | -0.005                | 2/9    | 0.180        | <b>+0.033</b>        | 12/12  | <b>0.000</b> | <b>+0.034</b> | 12/12  | <b>0.000</b> |
| Single-Truth Recall    | -0.005                | 3/9    | 0.508        | +0.026               | 9/12   | 0.146        | +0.025        | 8/12   | 0.388        |
| Decision accuracy      | -0.003                | 2/9    | 0.180        | -0.002               | 5/12   | 0.774        | -0.004        | 5/12   | 0.774        |

Table 19: **Mechanism-matched oracle ablations (fine-tuned models)**. Each intervention supplies gold information for one reasoning stage; a metric is *targeted* if it depends on that stage. The gold conflict label moves Behaviour Adherence and nothing else; the gold per-document notes move every grounding and citation metric in 9/9 cells and leave the answer/refusal decision untouched.  $p$  is an exact two-sided sign test over model $\times$ prompt cells. Qwen2.5-32B is excluded from the conflict-oracle columns because of the scoring artifact documented in Appendix L; the +/ $n$  column reports the exact number of cells contributing to each test ( $n = 9$  for the conflict oracle,  $n = 12$  otherwise).




Figure 13: Behaviour across the prompt ladder. Solid lines are trace-tuned models, dashed lines untuned. Fine-tuned behaviour is essentially flat as the instruction is stripped away; untuned grounding and refusal collapse.

error in place, which locates that error in Stage-3 synthesis, in gold-annotation noise, or in judge strictness; we have not separated those three.

For baseline models every oracle *hurts*, and the conflict label is catastrophic for grounding ( $-0.238$ ; Figure 16). The same injected information is informative for a trace-tuned model and




Figure 14: Behaviour Adherence under the **strict** prompt, baseline against trace-tuned. Deltas are annotated; bold green marks a gain exceeding the  $\pm 0.034$  noise floor.

harmful for a baseline one, which is a caution against reading oracle numbers as upper bounds on what prompting alone can achieve.

## L Excluded Rows and Diagnostics

Three rows are excluded from all reporting: Qwen2.5-32B, conflict-oracle, all three prompts. Factual Grounding collapses to 0.0016–0.0047 while Behaviour Adherence ( $\approx 0.82$ ) and Single-Truth Recall ( $\approx 0.75$ ) stay normal, against FG of 0.81–0.88 for the same model elsewhere, a roughly two-hundred-fold drop isolated to one cell. Normal behaviour and recall with near-zero grounding means the model produced adequate, fact-containing answers whose claim-level citation link-






Figure 15: Behaviour Adherence under the **minimal** prompt. Baselines compress into a narrow band once the instruction no longer names the taxonomy, and every model gains.




Figure 16: Mean effect of each oracle on every metric. Fine-tuned models gain on the targeted metrics; untuned models are harmed uniformly, most severely on grounding under the conflict-label oracle.

age went unrecovered by the scorer; the independently computed citation metrics for these runs are ordinary, localising the fault to grounding extraction rather than generation. Because answer quality fuses geometrically with FG, these rows also yield the aggregate floor of exactly 0.400 (Appendix F.8). We exclude Qwen2.5-32B from both run types here so baseline and fine-tuned comparisons remain balanced.

Separately, three benchmark records carry a numeric conflict identifier disagreeing with their authoritative textual label; applicability from the numeric field would give a single-truth denominator of 484 rather than 487. Every reported run shows 487, so no result is affected.

## M Released Artifacts

We release the full resource and the prompt suite used to build and evaluate it. The release comprises: the **training/validation split** (943 examples) and the **held-out benchmark** (736 examples), together carrying 10,343 adjudicated evidence units, each with a per-document verdict, an extracted key fact, a snippet-anchored supporting quote, a verdict rationale, a source-quality label and the complete annotator vote tally; the **stagewise annotation prompts** (system prompt and user template for each of the three annotation stages), which define the annotation contract described in Appendix B; the **end-to-end inference prompts** for all three instruction strengths; the **CATS judge prompts**; and the per-run evaluation outputs from which every number in this paper is computed.

Records are distributed as JSONL with dataset cards documenting per-source provenance and the upstream licence for each source family. All artifacts are available at [https://anonymous.4open.science/r/rag\\_reason-2F8F/](https://anonymous.4open.science/r/rag_reason-2F8F/).

## N End-to-End Inference Prompts

### N.1 Why the prompt is varied

Varying the prompt is what makes it testable whether conflict-aware behaviour has been *learned* or is merely being *recited* from the instruction. A model that produces the trace only when the strict prompt describes it may be following the prompt; a model that still produces it under the minimal prompt, which never mentions the trace, the taxonomy or the citation format, has learned a mapping from retrieval context to behaviour. Training and evaluating across all three variants is therefore an instrument, not a data-augmentation trick, and §8.2 reports the resulting gap.

The three end-to-end instruction strengths are reproduced verbatim below. Each condition is a system prompt plus a user template; the placeholders are filled with the query and the retrieved documents at generation time. Only typographic normalisation was applied for inclusion here and no wording was changed.

The ladder is a controlled reduction in *instruction*, not in task. The underlying evidence problem, the retrieved set and the expected target are identical across the three; what changes is how much of the procedure the prompt restates (Table 20):

The minimal system prompt is four sentences long and *never mentions* the evidence trace, the




Figure 17: Each oracle moves the metric it targets and essentially nothing else. Green bars are targeted metrics, grey bars are specificity controls; diamonds mark means, points individual model × prompt cells.

#### Condition, length and what the prompt still states

|                |                                                                                                    |
|----------------|----------------------------------------------------------------------------------------------------|
| <b>strict</b>  | — 1,965 words (49 ×)                                                                               |
|                | full taxonomy, verdict heuristics, threshold and edge-case rules, citation policy, output contract |
| <b>runtime</b> | — 195 words (5 ×)                                                                                  |
|                | three-stage schema, allowed labels, abstention and composition rules, citations, sentinel          |
| <b>minimal</b> | — 40 words (1 ×)                                                                                   |
|                | evidence-only answering and when to abstain. Nothing else                                          |

Table 20: The three end-to-end instruction strengths, measured by system-prompt length.

per-document notes, or both. Those twelve further templates are released with the artifacts but are not reproduced here, since they differ from the corresponding end-to-end prompt only by the injected block and the instruction to preserve it.

1942  
1943  
1944  
1945  
1946

three stages, the conflict taxonomy, the citation format or the terminating sentinel. A fine-tuned model nonetheless emits all of them, which is what §8.2 measures: behaviour under this prompt is the internalization probe, and the collapse of untuned grounding here (0.031–0.108) against fine-tuned grounding (0.792–0.857) is a comparison between models given *these forty words* and nothing more.

Two details in the strict prompt are worth flagging because they are referenced elsewhere. It requires citations on at least 80% of answer sentences, whereas the evaluator marks a citation pass at 75%; the specification and the scoring threshold are related but not identical, and we report both rather than eliding the gap. It also asks the model to order cited documents by the source-quality label it assigned in Stage 1, which is one route by which the URL-whitelist heuristic of Appendix B can influence answer composition.

The three oracle conditions of §8.4 use the same three instruction strengths with an additional block injecting the gold conflict label, the gold

1947

## N.2 Strict

1948

### Strict system prompt (`system_e2e.txt`)

```
You are a conflict-aware Retrieval-Augmented Generation assistant. You must produce a
STRICT PUBLIC EVIDENCE TRACE followed by a grounded final answer, and comply with all
requirements below.

You will be given:
- a query, and
- retrieved documents (d1...dN) containing snippets and metadata.

You must use ONLY the provided retrieved snippets as evidence. You must not use external
knowledge or your own parametric knowledge.

#####
EVIDENCE ANCHORING FOR PER-DOC VERDICTS
#####

Primary evidence is the document's snippet in retrieved_docs.

key_fact:
- ONE sentence strictly entailed by the snippet.
- Every concrete value in key_fact (names/dates/locations/numbers) must appear in the
  snippet.

verdict_reason (brief):
- Must justify the verdict using only what the snippet states.
- Do NOT add new facts, sources, or interpretations beyond what the snippet states.

If you cannot clearly anchor a key_fact/verdict_reason to the snippet:
- Do NOT choose "supports".
- Choose "partially supports" if the doc is on-topic but incomplete/hedged/indirect.
- Choose "irrelevant" if it does not help answer the query, or it is irrelevant to the
  query.

Do not correct the snippet. Judge only what is written. Do not import external facts or use
your parametric knowledge.

#####
VERDICT HEURISTICS & THRESHOLDS
#####

- supports: the snippet directly answers the query with an explicit, decisive claim.
- partially supports: on-topic but incomplete, hedged, indirect, or missing a required
  detail.
- irrelevant: not useful for answering the query.

1) Threshold queries:
- For "over/at least X?": a span stating a maximum/ceiling/"at most"/a range whose upper
  bound <= X can be decisive.
- Categorical statements ("cannot exceed X") can be supports if clearly stated.
- Hedged language ("may/might/probably/likely") alone -> partially supports unless a
  decisive bound is stated.

2) Span preference:
- If multiple candidate facts exist, choose the most specific snippet content containing
  decisive values/dates/names.

3) Missing required details:
- If the query requires a date/number/name and the snippet lacks it: do not mark supports.
  Use partially supports if on-topic, else irrelevant.

4) "Next/most recent/latest":
- supports only if the snippet explicitly identifies the earliest/"next"/"most
  recent/latest".
- Lists without a clear "next/most recent/latest" -> partially supports.
- Mere description without "latest/next/most recent" language -> partially supports.

5) Comparisons/opinions:
- supports if a clear overall claim answers the general case.
- partially supports if limited to a subset/region/industry/conditional case or small
  samples.

6) Negative/inconclusive evidence:
- "No evidence / not enough evidence / inconclusive" -> partially supports (never
  supports).

7) Date-specific queries:
- supports only with a full calendar date or complete date range.
- Year-only/vague/contradictory dates -> partially supports.

8) Factual identification ("Who/What/Where/When"):
- supports if the snippet names the entity and the required status
  (current/latest/official) clearly.
```

1949

- Otherwise partially supports.

9) Exact count / tally questions:

- If the retrieved snippets explicitly enumerate the relevant entities or give an exact count, answer with that exact count.
- If an exact count can be derived directly from the retrieved snippets without outside knowledge, derive it carefully and answer with the exact count.
- If the snippets show only examples, approximations, ranges, or partial lists, do not invent a precise count; say the exact count cannot be determined from the retrieved evidence.

10) Ambiguous scope:

- Do not silently assume a country, region, organization, or time scope unless the query or evidence clearly makes that scope primary.
- If the evidence points to more than one plausible scope, state the scope explicitly in the first sentence and answer within that scope.
- If the query is broad and the evidence is mixed by jurisdiction or definition, do not overgeneralize from one country-specific source.

#####  
SOURCE PREFERENCE & CITATION POLICY  
#####

Assign source\_quality using the source\_url/source metadata when available.

High-credibility (prefer and cite first):

- official government or institutional sources (.gov, .edu)
- WHO/UN/CDC and official international organizations
- peer-reviewed journals and reputable encyclopedias (e.g., Britannica)
- major outlets (Reuters, BBC, AP, NYT, WSJ, Guardian)

Low-credibility:

- blogs, unverified forums, marketing pages, social media, miscellaneous sites

When multiple sources support the same fact, cite high-credibility sources first, then add others by decreasing utility.

#####  
CONFLICT TYPE DEFINITIONS  
#####

Choose exactly one label for the set of non-irrelevant documents:

1) No conflict

- The search results refer to the same concept and are in agreement. Differences are superficial, such as variations in detail or granularity.
- All non-irrelevant docs agree on the key claim; differences are superficial.

2) Complementary information

- The question is underspecified or allows for multiple valid perspectives or answers that do not contradict each other. All provided answers can be considered correct.
- Non-irrelevant docs add different facets (time, region, subgroup, definition) that fit together without contradiction.

3) Conflicting opinions or research outcomes

- The query addresses a subjective or contentious topic, leading to genuinely opposing viewpoints or contradictory research findings.
- Some non-irrelevant docs contradict each other within the same scope and time window; mutually exclusive claims or incompatible outcomes.
- Non-irrelevant docs are incompatible within the same scope and time window; mutually exclusive claims or incompatible outcomes.

4) Conflict due to outdated information

- The question has a verifiable factual answer, but the search results contain conflicting information due to changes over time.
- Newer or more up-to-date statements contradict or supersede older factual claims.

5) Conflict due to misinformation

- The question has a verifiable factual answer, but some of the search results contain factually incorrect or misleading information.
- Some sources are factually incorrect or misleading versus more reliable references within the retrieved set.
- Some retrieved claims are factually incorrect or misleading relative to other, more credible retrieved evidence.
- Use this label sparingly; choose it only when the retrieved set itself provides a clear basis to call something false or misleading with respect to the query.

#####  
ABSTENTION POLICY (STRICT)  
#####

Abstain only if ALL docs are "irrelevant" OR the set collectively fails to address the query.

If ANY doc has verdict "supports", do NOT abstain; produce the best supported answer with conflict-aware framing.



If all docs are "partially supports" or "irrelevant" (none "supports"), use judgment: abstain only if, even taken together, the partial evidence cannot form a meaningful response to the query; if together they are sufficient, do NOT abstain. Disagreement alone is NOT a reason to abstain. If the evidence supports a comparison or conflict-aware answer, answer it. Abstention is about inability to answer the query from the retrieved evidence, not about the presence of nuance or disagreement.

#####  
COMMON CONFUSION GUARDS  
#####

- Do NOT choose Complementary information when the retrieved docs make mutually incompatible claims within the same scope and time window.
- If the disagreement is best explained by time change or superseded facts, choose Conflict due to outdated information instead of Complementary information.
- If the disagreement is within the same scope/time window and reflects opposing viewpoints or incompatible findings, choose Conflicting opinions or research outcomes.
- Use Conflict due to misinformation only when the retrieved set itself clearly supports calling one claim false or misleading relative to stronger evidence.
- If there is at least one supports document, do not fall back to abstention just because some other documents are partial, noisy, or lower quality.

#####  
EXPECTED BEHAVIOR RULES (CONDITIONAL ON LABEL)  
#####

Your final answer must follow the chosen conflict type label:

- Conflict due to outdated information:  
Prioritize the most recent and credible information; acknowledge older or superseded claims when relevant.
- Conflicting opinions or research outcomes:  
Present differing perspectives neutrally, attribute differences to sources, and avoid taking sides unless the evidence clearly favors one.
- Conflict due to misinformation:  
Identify and correct false or unreliable claims using more credible documents within the retrieved set.
- Complementary information:  
Combine partial, non-contradictory facts into a coherent answer.
- No conflict:  
Answer directly using the strongest consistent evidence.

#####  
ANTI-FAILURE GUARDS  
#####

- Exactly one <think>...</think>.
- The literal string "<think>" must not appear inside the think block content.
- Stage 1 must enumerate dl...dN exactly once, in order, without gaps, fabrications, reordering, or ranges.
- Do not write doc-id ranges like "dl-d5" anywhere.
- Stage 2 must choose exactly one allowed conflict label.
- Use the exact conflict label strings listed above. Do not invent labels.
- At least 80% of final-answer sentences include [dX] unless abstaining.
- A non-abstain final answer with zero [dX] citations is invalid.
- Use only in-range citations; never cite out-of-bounds.
- No new facts; adhere to length budgets.
- Be brief in the prose sections of the think block. Extra explanation that is not needed increases failure risk.
- End with [[END-OF-ANSWER]].

FINAL SELF-CHECK BEFORE YOU FINISH

- Before emitting the final token, verify that you have written exactly one opening "<think>" line and one closing "</think>" line.
- If you are abstaining, the final answer must be exactly: CANNOT ANSWER, INSUFFICIENT EVIDENCE
- If you are not abstaining, ensure at least 80% of final-answer sentences end with one or more [dX] citations.
- If the final line is not yet [[END-OF-ANSWER]], write [[END-OF-ANSWER]] exactly as the last line.

#####  
OUTPUT CONTRACT (STRICT 3-STAGE TRACE)  
#####

You must output EXACTLY in this order, with no extra text before or after:

- 1) A line that is exactly:  
<think>

2) Inside the think block, write these three stages in this exact order:

Stage 1 - Evidence assessment:

- d1: supports|partially supports|irrelevant - brief verdict\_reason grounded only in d1.  
Key fact: brief snippet-entailed key fact, or empty if irrelevant. Source quality: high|low.
- d2: supports|partially supports|irrelevant - brief verdict\_reason grounded only in d2.  
Key fact: brief snippet-entailed key fact, or empty if irrelevant. Source quality: high|low.
- Continue until EVERY retrieved doc d1...dN is covered exactly once, in order.

Stage 1 requirements:

- Use only these verdict labels: supports, partially supports, irrelevant.
- If verdict is irrelevant, the key fact must be empty or explicitly state that no useful key fact is present.
- Do not fabricate, skip, merge, reorder, or range-compress doc ids.

Stage 2 - Conflict assessment:

Conflict type: <one allowed conflict label>

Reason: <concise conflict\_reason, <=50 words>

Evidence pattern: <brief synthesis of how the documents agree, complement, conflict, are outdated, contain misinformation, or fail to support an answer>

Stage 2 requirements:

- The conflict type MUST be exactly one of:  
No conflict  
Complementary information  
Conflicting opinions or research outcomes  
Conflict due to outdated information  
Conflict due to misinformation
- Keep the label consistent with the evidence pattern and final answer behavior.

Stage 3 - Answer plan:

<1-2 concise sentences explaining how the cited evidence yields the final answer, or why the evidence requires abstention.>

3) A line that is exactly:  
</think>

4) ONE BLANK LINE.

5) The FINAL ANSWER lines.

6) A line that is exactly:  
[[END-OF-ANSWER]]

FINAL ANSWER REQUIREMENTS

- If you are abstaining: the final answer line must be EXACTLY:  
CANNOT ANSWER, INSUFFICIENT EVIDENCE  
No citations must appear in an abstain answer.
- Otherwise: write 2-4 sentences, or 4-5 for simple unanimous facts.
  - Use bracketed citations [dX]; at least 80% of sentences MUST include at least one citation of the form [dX] at the end of that sentence.
  - Cite only existing doc\_ids (d1...dN); never cite [dK] where K does not belong to the retrieved set.
  - Prefer ordering of cited docs by their credibility (docs you labeled source\_quality="high" first, then "low"), then by utility.
  - For exact-count questions, prefer a precise count when the retrieved evidence directly supports it; otherwise explicitly say the exact count cannot be determined from the retrieved evidence.
  - For broad or ambiguous queries, name the scope you are answering within instead of implying a universal answer that the snippets do not establish.

No markdown fences, no extra headings, no extra commentary anywhere.  
There must be EXACTLY ONE <think>...</think> block (no nesting, no repeats).  
The response MUST end with the sentinel line [[END-OF-ANSWER]].

1952

## Strict user template (user\_e2e.txt)

1953

You are given a query and retrieved documents.  
Use only the retrieved snippets as evidence.

Query:  
{query}

Retrieved documents (JSON; 'doc\_id' is canonical and must be preserved exactly in your evidence trace and citations):  
{retrieved\_docs\_json}

1954

## N.3 Runtime

1955

## Runtime system prompt (system\_e2e\_runtime.txt)

1956

You are a conflict-aware RAG assistant.

Use only the retrieved documents. Do not use outside knowledge.

Output a public 3-stage evidence trace followed by the final answer:

```
<think>
Stage 1 - Evidence assessment:
- dl: supports|partially supports|irrelevant - brief reason grounded in the snippet.

Stage 2 - Conflict assessment:
Conflict type: one allowed conflict label
Reason: concise reason
Evidence pattern: brief synthesis of how the documents agree, complement, conflict, or fail.

Stage 3 - Answer plan:
Brief plan for the final answer or grounded refusal.
</think>
```

Then write the final answer with end-of-sentence [dX] citations, or exactly:  
CANNOT ANSWER, INSUFFICIENT EVIDENCE

End with [[END-OF-ANSWER]].

Allowed conflict labels: No conflict; Complementary information; Conflicting opinions or research outcomes; Conflict due to outdated information; Conflict due to misinformation.

Abstain only if all documents are irrelevant or the retrieved set still cannot answer the query.  
If any document supports the answer, do not abstain.  
If no single document fully answers the query but multiple partially supporting documents jointly determine the answer, do not abstain; synthesize them carefully.  
Disagreement alone is not a reason to abstain if a grounded conflict-aware answer can still be given.  
Cite only retrieved document ids.

#### Runtime user template (user\_e2e\_runtime.txt)

```
Query:
{query}

Retrieved documents:
{retrieved_docs_json}
```

## N.4 Minimal

#### Minimal system prompt (system\_e2e\_minimal.txt)

You are a RAG assistant.

Answer the query using only the retrieved documents.  
If the documents jointly answer the query, answer it even when each one is only partial.  
Abstain only when the retrieved evidence still leaves a necessary gap.

#### Minimal user template (user\_e2e\_minimal.txt)

```
Query:
{query}

Retrieved documents:
{retrieved_docs_json}
```

## O CATS Judge Prompts

The complete active judge bundle is reproduced verbatim below: one shared system message, the three task prompts, and the five-type behaviour rubric that the Behaviour Adherence prompt selects from. Placeholders in braces are filled at evaluation time. Only typographic normalisation was applied.

This is the *entire* bundle. Grounded Refusal has no prompt because it is computed deterministically, and no natural-language-inference prompt appears because the active grounding path does not use one; earlier NLI-based grounding code exists in the repository but is not the protocol behind any number reported here.

Three properties of these prompts bear on claims made in the main text, and two of them are qualifications rather than confirmations.

**Behaviour is explicitly walled off from the other components.** The Behaviour Adherence prompt instructs the judge to ignore factual correctness and states outright that answerability correctness, factual entailment, citation validity, unsupported-claim detection and single-truth recall must not be used as hidden criteria. It supplies exactly one rubric line, for the sample’s own conflict type, rather than asking the judge to score all five behaviours. This is the orthogonality requirement of §6 in its operational form, and it is corroborated empirically: Behaviour Adherence is the CATS component least correlated with citation behaviour (Appendix H). The prompt also defines a refusal as adherent when the evidence genuinely does not support an answer, which is why correct refusals are excluded from the behaviour denominator rather than scored as failures.

**Factual Grounding is deliberately permissive, so it is an upper bound on strict grounding.** The prompt enumerates seven distinct routes to support, including paraphrase, entailment, more-specific instances, more-general statements, containing ranges and partial coverage of a claim’s substance, and it instructs the judge to “bias toward inclusion when in doubt”. This is a design choice: the metric is meant to measure whether an answer *uses and cites* retrieved evidence, not whether each claim is strictly entailed. It does mean our grounding numbers should be read as generous rather than conservative, and that a stricter entailment-based variant would score lower.

**One wording inaccuracy we do not wish to paper over.** The grounding prompt tells the judge that the per-document verdicts it is shown are “ground truth annotations made by human experts”. They are in fact produced by the annotation committee of Appendix B. Human review covered the conflict-type label and the benchmark preselection fields, but *not* the per-document verdicts (Appendix C). The instruction’s operative function is to stop the judge re-litigating document relevance, which it does regardless of provenance, so we do not expect a material effect on the reported grounding scores; but the description is inaccurate and a corrected wording should be used in any future version.

## O.1 Shared system message

### Committee system message (`committee_json_system_prompt.txt`)

```
You are an evaluation judge. Respond ONLY with JSON.
```

## O.2 Behaviour Adherence

### Behaviour Adherence prompt (`behavior_adherence_prompt.template.txt`)

```
You are evaluating ONLY the behavior of a model answer, not its factual correctness.

Behavior means:
- How the answer handles multiple sources, uncertainty, disagreement, or lack of conflict.
- Whether it summarizes, reconciles, or contrasts viewpoints as appropriate.
- Whether it is direct vs. hedged, neutral vs. biased, etc.

Given:
- A user query
- A model-generated answer
- A conflict type with ONE expected behavior rubric (shown below)

Your task:
Decide whether the model's answer follows the expected behavior for THIS specific
conflict type only.

Conflict Type: {conflict_type}
Expected Behavior for this conflict type:
{expected_behavior}
{provenance_block}
Instructions:
- Judge the answer ONLY against the rubric shown above. Do not invoke rubrics for other
  conflict types.
- If the answer is a refusal ("CANNOT ANSWER" / "INSUFFICIENT EVIDENCE") AND the retrieved
  evidence genuinely does not support an answer, treat the refusal as adherent.
- If the answer clearly follows the expected behavior for this conflict type, set
```



```

    "adherent": true.
- If the answer clearly violates or ignores the expected behavior, set "adherent": false.
- Do not use answerability correctness, factual entailment, citation validity,
  unsupported-claim detection, or single-truth recall as hidden behavior criteria. Those
  are separate metrics.
- Ignore factual correctness; only judge how the answer behaves relative to the rubric.
- The "rationale" should briefly cite the key aspects of the answer's behavior.
- The "confidence" should reflect how certain you are in YOUR adherent decision, between
  0.0 (coin flip) and 1.0 (clearly correct call). Use 0.5 if you are uncertain.

Return ONLY a JSON object with fields:
  "adherent": true or false,
  "rationale": "short explanation",
  "confidence": 0.0 to 1.0

Query:
{query}

Answer:
{answer}

```

## Five-type behaviour rubric (behavior\_rubric.md)

```

# Behavior Adherence Rubric

This is the exact five-type rubric selected by
`behavior_judge_prompt(...)` in `rag_eval/judge_prompts.py`.

| Conflict type | Rubric supplied to the judge |
| --- | --- |
| 1. No Conflict | Provide a clear and direct answer without introducing alternative
  viewpoints or uncertainty. |
| 2. Complementary Information | Consolidate and reconcile the different partial answers
  provided by the retrieved documents, without framing the response as a debate. |
| 3. Conflicting Opinions or Research Outcomes | Explicitly reflect the debate between the
  retrieved sources and neutrally summarize the different viewpoints. |
| 4. Conflict Due to Outdated Information | Prioritise the up-to-date information, while
  optionally acknowledging the presence of outdated sources. |
| 5. Conflict Due to Misinformation | Disregard inaccurate sources and provide a response
  grounded in reliable and verified information. |

The active prompt selects exactly one rubric for the supplied conflict type. It
does not ask the judge to score all five behaviors at once.

The BA judge is instructed to assess behavior only. It must not use answerability
correctness, factual entailment, citation validity, unsupported-claim detection,
or Single-Truth Recall as hidden BA criteria. A refusal is adherent when the
retrieved evidence genuinely does not support a confident answer, regardless of
conflict type.

```

## 0.3 Factual Grounding

### Factual Grounding prompt (factual\_grounding\_prompt.template.txt)

```

You are a fact-checking judge. Your task is to determine which retrieved documents support
a specific claim from a model's answer.

IMPORTANT GROUND RULES:
- The document verdicts ('supports', 'partially supports') are GROUND TRUTH annotations
  made by human experts. You do NOT need to re-verify whether these docs are relevant to
  the query , they already are. Your job is ONLY to check whether the SPECIFIC CLAIM is
  conveyed by each doc's text.
- Do NOT use outside knowledge. Only what the document text says.
- No concept of 'contradicts' , ignore any conflict between docs. Only look for support.

WHAT COUNTS AS SUPPORT (be generous , partial and paraphrastic support BOTH count):
A document supports the claim if its text establishes the claim's core assertion. The
wording does NOT need to match the claim word-for-word. In particular, the doc supports
the claim when ANY of the following holds:
(a) The doc paraphrases the claim, uses synonyms, or summarizes the same fact.
(b) The doc states a fact that, taken at face value, entails or implies the claim.
(c) The doc gives a more specific instance, example, or sub-case of the claim (specific
  instance -> the general claim is supported).
(d) The doc gives a more general statement that covers the claim (the claim is a natural
  sub-case of what the doc says).
(e) The doc supplies a numeric range, qualitative range, or category that includes the
  value the claim asserts.
(f) The doc covers a meaningful portion of the claim's substance , even if secondary
  modifiers, qualifiers, or peripheral details in the claim aren't in the doc. Missing
  modifiers do NOT disqualify a doc; missing the core fact does.

```

(g) The gold verdict is 'partially supports' AND the doc's key\_fact/passage addresses the claim's main assertion , treat this as supporting unless the doc is clearly about a different topic.

WHAT DOES NOT COUNT AS SUPPORT:

- The doc is about a different entity, topic, or time period than the claim.
- The doc only mentions a keyword from the claim without addressing the claim's actual assertion.
- The doc would require external knowledge or several inferential leaps to connect to the claim.

Bias toward INCLUSION when in doubt: if a reasonable reader would say 'yes, this doc backs up the gist of the claim,' mark it supporting. Reserve exclusion for docs that are clearly off-topic or only tangentially related.

QUERY: {query}

{model\_answer\_block}SPECIFIC CLAIM TO EVALUATE: "{claim\_text}"

RETRIEVED DOCUMENTS (pre-annotated supporting docs , decide which contain evidence for the SPECIFIC CLAIM above):  
{documents\_block}

YOUR TASKS:

1. For each document above, apply the WHAT COUNTS AS SUPPORT rules (a)-(g). List every doc that meets ANY of (a)-(g) in supporting\_docs.
2. If NO single document satisfies (a)-(g), check whether the COMBINATION of any two documents together supports the claim (cross-document evidence , one doc provides part A, another provides part B, and together they establish the claim). Only set cross\_doc\_support=true in this case.

Valid document IDs you may reference: {valid\_document\_ids}

Return ONLY a valid JSON object with exactly these keys and no comments, no markdown, and no extra text:

```
{
  "supporting_docs": ["d2", "d3"],
  "cross_doc_support": false,
  "cross_doc_combo": []
}
```

Rules for the JSON fields:

- "supporting\_docs": doc\_ids that individually support the claim; use [] if none.
- "cross\_doc\_support": true ONLY when no single doc supports but two docs together do.
- "cross\_doc\_combo": exactly the two doc\_ids that combine to support when cross\_doc\_support is true; otherwise [].

2010

## O.4 Single-Truth Recall

2011

### Single-Truth Recall prompt (single\_truth\_recall\_prompt.template.txt)

2012

You are checking whether a model answer correctly **asserts** a given factual answer.

The gold answer is the correct answer to a question. Your job is to decide whether the model's answer

**commits to** (asserts) this gold answer as the model's own conclusion.

MATCHING RULES (count as adherent=true):

- Paraphrases, abbreviations, and equivalent wording that convey the same fact.
- **Logically equivalent statements** (e.g. "born in 1990" vs "born thirty-five years ago in 2025";  
"A is the capital of B" vs "B's capital is A").
- Minor misspellings, casing differences, whitespace, punctuation, or unit-formatting differences  
(e.g. "Chiliwack" ~="Chilliwack"; "Stephan Curry" ~="Stephen Curry"; "1,759" ~="1759"; "\$1.8 billion" ~="1.8B USD").

NON-MATCHING RULES (count as adherent=false):

- The model commits to a DIFFERENT answer, even if the gold answer also appears somewhere in the text  
as a **quotation, citation, or attribution** to a document.  
Example: gold="1,759"; model says "Document d2 reports 1,759, but the correct figure is 658." ->  
adherent=false (model asserts 658, only mentions 1,759 as a cited claim it rejects).
- The model only lists the gold answer as one of several possibilities without endorsing it.
- The model refuses ("CANNOT ANSWER" / "INSUFFICIENT EVIDENCE") , recall is 0 regardless of what documents say.
- The model asserts something logically incompatible with the gold answer.

Return ONLY a JSON object with fields:

"adherent": true or false (whether the model's answer ASSERTS the gold answer as its own conclusion),  
"rationale": short string explanation (briefly note whether the gold was asserted or

2013

```

    merely mentioned),
    "confidence": 0.0 to 1.0 (your certainty in the adherent decision)

Gold factual answer:
{gold_answer}

Model answer:
{model_answer}

```

## P Complete Agreement Statistics

The tables in this section report *every* agreement statistic computed in this work, so that no coefficient in the main text is quoted without its denominator, its stratum and its confusion counts. Four conventions apply throughout. Raw agreement and  $\kappa$  answer different questions and are always reported together, because a strongly imbalanced field can show high raw agreement and near-zero  $\kappa$  without either number being wrong. A dash in a  $\kappa$  column marks a field that is single-valued across its stratum, where chance-corrected agreement is undefined; we report exact agreement rather than substituting a coefficient. Nothing is imputed for missing judgments. And committee-internal agreement is agreement among model judges under a fixed prompt and serving protocol, never to be read as human inter-annotator agreement.

Table 21: **Complete resource-review agreement statistics.** Every recorded field in both review workflows.  $n$  is the number of review pairs; *agree/disagree* are exact counts. A dash in the  $\kappa$  column marks a field that is single-valued across the stratum, where chance-corrected agreement is undefined and exact agreement is the only meaningful statistic.

| Workflow / scope                                                            | Field                      | $n$  | agree | dis. | raw   | $\kappa$ |
|-----------------------------------------------------------------------------|----------------------------|------|-------|------|-------|----------|
| <i>Training conflict-type review (943 records, 3 reviewers, 2 per item)</i> |                            |      |       |      |       |          |
|                                                                             | reviewed_conflict_type     | 943  | 787   | 156  | 0.835 | 0.7694   |
|                                                                             | label_action               | 943  | 791   | 152  | 0.839 | 0.1341   |
|                                                                             | changed_label              | 943  | 791   | 152  | 0.839 | 0.1341   |
|                                                                             | review_confidence          | 943  | 607   | 336  | 0.644 | 0.0159   |
| <i>Benchmark preselection (1,454 records, 7 reviewers)</i>                  |                            |      |       |      |       |          |
|                                                                             | human_preselect_decision   | 1454 | 1428  | 26   | 0.982 | 0.9228   |
|                                                                             | preliminary_conflict_type  | 1454 | 1378  | 76   | 0.948 | 0.9217   |
|                                                                             | preselection_confidence    | 1454 | 1445  | 9    | 0.994 | 0.9767   |
|                                                                             | retrieval_quality          | 1454 | 1436  | 18   | 0.988 | 0.9550   |
|                                                                             | evidence_sufficiency       | 1454 | 1436  | 18   | 0.988 | 0.9600   |
|                                                                             | conflict_clarity           | 1454 | 1441  | 13   | 0.991 | 0.9660   |
|                                                                             | query_specificity          | 1454 | 1449  | 5    | 0.997 | 0.9875   |
|                                                                             | source_reliability         | 1454 | 1446  | 8    | 0.994 | 0.9881   |
|                                                                             | relevant_doc_count_bin     | 1454 | 1426  | 28   | 0.981 | 0.9527   |
|                                                                             | gold_answer_possible       | 1454 | 1440  | 14   | 0.990 | 0.9669   |
| <i>Refusal-quality stratum (128 released refusal items)</i>                 |                            |      |       |      |       |          |
|                                                                             | human_preselect_decision   | 128  | 128   | 0    | 1.000 | –        |
|                                                                             | preliminary_conflict_type  | 128  | 128   | 0    | 1.000 | 1.0000   |
|                                                                             | preselection_confidence    | 128  | 128   | 0    | 1.000 | –        |
|                                                                             | retrieval_quality          | 128  | 128   | 0    | 1.000 | –        |
|                                                                             | evidence_sufficiency       | 128  | 128   | 0    | 1.000 | –        |
|                                                                             | conflict_clarity           | 128  | 128   | 0    | 1.000 | –        |
|                                                                             | query_specificity          | 128  | 128   | 0    | 1.000 | –        |
|                                                                             | source_reliability         | 128  | 128   | 0    | 1.000 | –        |
|                                                                             | relevant_doc_count_bin     | 128  | 128   | 0    | 1.000 | –        |
|                                                                             | gold_answer_possible       | 128  | 128   | 0    | 1.000 | –        |
|                                                                             | refusal_required           | 128  | 128   | 0    | 1.000 | –        |
|                                                                             | refusal_ground_truth_valid | 128  | 128   | 0    | 1.000 | –        |
|                                                                             | refusal_rationale_quality  | 128  | 128   | 0    | 1.000 | –        |
|                                                                             | refusal_quality_label      | 128  | 128   | 0    | 1.000 | –        |

Table 21 covers both resource-review workflows. Two features are worth naming. In the training workflow, the retain-or-change and confidence fields show high raw agreement with very low  $\kappa$ ; this is a prevalence effect, since most reviewers accepted the committee label and reported high confidence, and it is why the five-way taxonomy field is the appropriate headline statistic rather than an average across

fields. In the refusal-quality stratum, every field is unanimous across all 128 items, which is why those rows carry exact agreement and no coefficient.

Reviewers are denoted R1–R4 throughout; the mapping is fixed across all tables in this appendix but is otherwise arbitrary and carries no ordering or ranking.

Table 22: **Complete CATS human-evaluation binary agreement.** All strata for all four comparison frames.  $P_{pos}/P_{neg}$  are positive- and negative-specific agreement;  $n_{11}$ – $n_{00}$  are the confusion cells;  $\alpha$  is nominal Krippendorff’s alpha where computed. Conflicting opinions (type 3) is excluded from Single-Truth Recall by construction and so has no STR stratum. Human–human rows are response-level over the 300-response frame; human–committee rows are judgment-level and therefore pool all accepted reviews.

| Frame                        | Metric     | Stratum                | $n$  | agr   | $\kappa$ | $P_{pos}$ | $P_{neg}$ | $n_{11}$ | $n_{10}$ | $n_{01}$ | $n_{00}$ | $\alpha$ |
|------------------------------|------------|------------------------|------|-------|----------|-----------|-----------|----------|----------|----------|----------|----------|
| Human vs human               | Behaviour  | overall                | 300  | 0.853 | 0.669    | 0.891     | 0.776     | 180      | 33       | 11       | 76       | 0.667    |
|                              | Behaviour  | R1 / R2                | 83   | 0.819 | 0.587    | 0.870     | 0.706     | 50       | 14       | 1        | 18       | 0.578    |
|                              | Behaviour  | R1 / R3                | 84   | 0.821 | 0.595    | 0.870     | 0.717     | 50       | 13       | 2        | 19       | 0.589    |
|                              | Behaviour  | R1 / R4                | 21   | 0.952 | 0.859    | 0.970     | 0.889     | 16       | 0        | 1        | 4        | 0.862    |
|                              | Behaviour  | R2 / R3                | 83   | 0.867 | 0.719    | 0.893     | 0.825     | 46       | 5        | 6        | 26       | 0.720    |
|                              | Behaviour  | R2 / R4                | 19   | 0.895 | 0.756    | 0.923     | 0.833     | 12       | 1        | 1        | 5        | 0.763    |
|                              | Behaviour  | R3 / R4                | 10   | 1.000 | 1.000    | 1.000     | 1.000     | 6        | 0        | 0        | 4        | 1.000    |
|                              | Behaviour  | type 1 No conflict     | 63   | 0.937 | 0.714    | 0.964     | 0.750     | 53       | 2        | 2        | 6        | 0.716    |
|                              | Behaviour  | type 2 Complementary   | 60   | 0.783 | 0.536    | 0.835     | 0.683     | 33       | 12       | 1        | 14       | 0.522    |
|                              | Behaviour  | type 3 Conflicting op. | 58   | 0.862 | 0.725    | 0.862     | 0.862     | 25       | 6        | 2        | 25       | 0.727    |
|                              | Behaviour  | type 4 Outdated        | 60   | 0.817 | 0.505    | 0.879     | 0.621     | 40       | 8        | 3        | 9        | 0.504    |
|                              | Behaviour  | type 5 Misinformation  | 59   | 0.864 | 0.725    | 0.879     | 0.846     | 29       | 5        | 3        | 22       | 0.727    |
|                              | STR        | overall                | 215  | 0.902 | 0.773    | 0.929     | 0.842     | 138      | 19       | 2        | 56       | 0.772    |
|                              | STR        | R1 / R2                | 61   | 0.902 | 0.783    | 0.925     | 0.857     | 37       | 5        | 1        | 18       | 0.784    |
|                              | STR        | R1 / R3                | 57   | 0.860 | 0.672    | 0.900     | 0.765     | 36       | 8        | 0        | 13       | 0.668    |
|                              | STR        | R1 / R4                | 16   | 0.938 | 0.818    | 0.960     | 0.857     | 12       | 1        | 0        | 3        | 0.823    |
|                              | STR        | R2 / R3                | 61   | 0.902 | 0.765    | 0.930     | 0.833     | 40       | 5        | 1        | 15       | 0.766    |
|                              | STR        | R2 / R4                | 12   | 1.000 | 1.000    | 1.000     | 1.000     | 8        | 0        | 0        | 4        | 1.000    |
|                              | STR        | R3 / R4                | 8    | 1.000 | 1.000    | 1.000     | 1.000     | 5        | 0        | 0        | 3        | 1.000    |
|                              | STR        | type 1 No conflict     | 56   | 0.911 | 0.711    | 0.945     | 0.762     | 43       | 5        | 0        | 8        | 0.710    |
|                              | STR        | type 2 Complementary   | 43   | 0.814 | 0.549    | 0.871     | 0.667     | 27       | 7        | 1        | 8        | 0.543    |
|                              | STR        | type 4 Outdated        | 57   | 0.930 | 0.812    | 0.953     | 0.857     | 41       | 4        | 0        | 12       | 0.812    |
|                              | STR        | type 5 Misinformation  | 59   | 0.932 | 0.865    | 0.931     | 0.933     | 27       | 3        | 1        | 28       | 0.866    |
|                              | FG claim   | overall                | 477  | 0.897 | 0.780    | 0.918     | 0.861     | 276      | 25       | 24       | 152      | 0.780    |
|                              | FG claim   | R1 / R2                | 137  | 0.905 | 0.786    | 0.930     | 0.854     | 86       | 0        | 13       | 38       | 0.784    |
|                              | FG claim   | R1 / R3                | 127  | 0.945 | 0.888    | 0.952     | 0.936     | 69       | 4        | 3        | 51       | 0.888    |
|                              | FG claim   | R1 / R4                | 35   | 0.914 | 0.821    | 0.930     | 0.889     | 20       | 0        | 3        | 12       | 0.822    |
|                              | FG claim   | R2 / R3                | 131  | 0.840 | 0.663    | 0.874     | 0.779     | 73       | 21       | 0        | 37       | 0.655    |
|                              | FG claim   | R2 / R4                | 27   | 1.000 | 1.000    | 1.000     | 1.000     | 19       | 0        | 0        | 8        | 1.000    |
|                              | FG claim   | R3 / R4                | 20   | 0.750 | 0.519    | 0.783     | 0.706     | 9        | 0        | 5        | 6        | 0.501    |
|                              | FG claim   | type 1 No conflict     | 77   | 0.987 | 0.972    | 0.990     | 0.982     | 49       | 0        | 1        | 27       | 0.972    |
|                              | FG claim   | type 2 Complementary   | 109  | 0.908 | 0.803    | 0.928     | 0.875     | 64       | 6        | 4        | 35       | 0.803    |
|                              | FG claim   | type 3 Conflicting op. | 115  | 0.887 | 0.749    | 0.914     | 0.835     | 69       | 6        | 7        | 33       | 0.750    |
|                              | FG claim   | type 4 Outdated        | 85   | 0.800 | 0.566    | 0.844     | 0.721     | 46       | 10       | 7        | 22       | 0.568    |
|                              | FG claim   | type 5 Misinformation  | 91   | 0.912 | 0.821    | 0.923     | 0.897     | 48       | 3        | 5        | 35       | 0.821    |
| Human vs committee           | Behaviour  | overall                | 650  | 0.745 | 0.407    | 0.814     | 0.591     | 364      | 68       | 98       | 120      | –        |
|                              | Behaviour  | type 1 No conflict     | 133  | 0.865 | 0.325    | 0.924     | 0.400     | 109      | 7        | 11       | 6        | –        |
|                              | Behaviour  | type 2 Complementary   | 130  | 0.708 | 0.279    | 0.804     | 0.424     | 78       | 4        | 34       | 14       | –        |
|                              | Behaviour  | type 3 Conflicting op. | 128  | 0.820 | 0.641    | 0.816     | 0.824     | 51       | 13       | 10       | 54       | –        |
|                              | Behaviour  | type 4 Outdated        | 130  | 0.746 | 0.338    | 0.829     | 0.507     | 80       | 19       | 14       | 17       | –        |
|                              | Behaviour  | type 5 Misinformation  | 129  | 0.581 | 0.149    | 0.630     | 0.518     | 46       | 25       | 29       | 29       | –        |
|                              | Behaviour  | llama8b                | 328  | 0.738 | 0.355    | 0.817     | 0.538     | 192      | 39       | 47       | 50       | –        |
|                              | Behaviour  | qwen7b                 | 322  | 0.752 | 0.451    | 0.811     | 0.636     | 172      | 29       | 51       | 70       | –        |
|                              | STR strict | overall                | 465  | 0.908 | 0.790    | 0.931     | 0.858     | 292      | 24       | 19       | 130      | –        |
|                              | STR strict | type 1 No conflict     | 118  | 0.924 | 0.769    | 0.952     | 0.816     | 89       | 8        | 1        | 20       | –        |
|                              | STR strict | type 2 Complementary   | 94   | 0.851 | 0.657    | 0.891     | 0.767     | 57       | 8        | 6        | 23       | –        |
|                              | STR strict | type 4 Outdated        | 124  | 0.935 | 0.841    | 0.955     | 0.886     | 85       | 6        | 2        | 31       | –        |
|                              | STR strict | type 5 Misinformation  | 129  | 0.907 | 0.814    | 0.910     | 0.903     | 61       | 2        | 10       | 56       | –        |
|                              | STR strict | llama8b                | 233  | 0.914 | 0.801    | 0.938     | 0.863     | 150      | 11       | 9        | 63       | –        |
|                              | STR strict | qwen7b                 | 232  | 0.901 | 0.779    | 0.925     | 0.854     | 142      | 13       | 10       | 67       | –        |
|                              | STR soft   | overall                | 465  | 0.912 | 0.798    | 0.935     | 0.863     | 295      | 21       | 20       | 129      | –        |
|                              | STR soft   | type 1 No conflict     | 118  | 0.924 | 0.762    | 0.952     | 0.809     | 90       | 7        | 2        | 19       | –        |
|                              | STR soft   | type 2 Complementary   | 94   | 0.872 | 0.701    | 0.908     | 0.793     | 59       | 6        | 6        | 23       | –        |
|                              | STR soft   | type 4 Outdated        | 124  | 0.935 | 0.841    | 0.955     | 0.886     | 85       | 6        | 2        | 31       | –        |
|                              | STR soft   | type 5 Misinformation  | 129  | 0.907 | 0.814    | 0.910     | 0.903     | 61       | 2        | 10       | 56       | –        |
|                              | STR soft   | llama8b                | 233  | 0.923 | 0.818    | 0.944     | 0.873     | 153      | 8        | 10       | 62       | –        |
|                              | STR soft   | qwen7b                 | 232  | 0.901 | 0.779    | 0.925     | 0.854     | 142      | 13       | 10       | 67       | –        |
|                              | FG claim   | overall                | 1029 | 0.899 | 0.783    | 0.920     | 0.862     | 599      | 38       | 66       | 326      | –        |
|                              | FG claim   | type 1 No conflict     | 165  | 0.945 | 0.880    | 0.958     | 0.922     | 103      | 4        | 5        | 53       | –        |
|                              | FG claim   | type 2 Complementary   | 233  | 0.910 | 0.805    | 0.929     | 0.876     | 138      | 8        | 13       | 74       | –        |
|                              | FG claim   | type 3 Conflicting op. | 250  | 0.888 | 0.759    | 0.911     | 0.848     | 144      | 12       | 16       | 78       | –        |
|                              | FG claim   | type 4 Outdated        | 183  | 0.825 | 0.606    | 0.870     | 0.733     | 107      | 9        | 23       | 44       | –        |
|                              | FG claim   | type 5 Misinformation  | 198  | 0.929 | 0.855    | 0.939     | 0.917     | 107      | 5        | 9        | 77       | –        |
|                              | FG claim   | llama8b                | 518  | 0.880 | 0.711    | 0.916     | 0.795     | 336      | 21       | 41       | 120      | –        |
|                              | FG claim   | qwen7b                 | 511  | 0.918 | 0.834    | 0.926     | 0.907     | 263      | 17       | 25       | 206      | –        |
| Human consensus vs committee | Behaviour  | overall                | 256  | 0.789 | 0.483    | 0.852     | 0.630     | 156      | 24       | 30       | 46       | –        |
|                              | Behaviour  | type 1 No conflict     | 59   | 0.881 | 0.299    | 0.935     | 0.364     | 50       | 3        | 4        | 2        | –        |

| Frame                     | Metric     | Stratum                       | $n$ | agr   | $\kappa$ | $P_{pos}$ | $P_{neg}$ | $n_{11}$ | $n_{10}$ | $n_{01}$ | $n_{00}$ | $\alpha$ |
|---------------------------|------------|-------------------------------|-----|-------|----------|-----------|-----------|----------|----------|----------|----------|----------|
|                           | Behaviour  | type 2 Complementary          | 47  | 0.787 | 0.391    | 0.865     | 0.500     | 32       | 1        | 9        | 5        | –        |
|                           | Behaviour  | type 3 Conflicting op.        | 50  | 0.880 | 0.760    | 0.885     | 0.875     | 23       | 2        | 4        | 21       | –        |
|                           | Behaviour  | type 4 Outdated               | 49  | 0.796 | 0.419    | 0.868     | 0.545     | 33       | 7        | 3        | 6        | –        |
|                           | Behaviour  | type 5 Misinformation         | 51  | 0.588 | 0.165    | 0.632     | 0.533     | 18       | 11       | 10       | 12       | –        |
|                           | STR strict | overall                       | 194 | 0.948 | 0.876    | 0.964     | 0.912     | 132      | 6        | 4        | 52       | –        |
|                           | STR strict | type 1 No conflict            | 51  | 0.961 | 0.865    | 0.976     | 0.889     | 41       | 2        | 0        | 8        | –        |
|                           | STR strict | type 2 Complementary          | 35  | 0.914 | 0.767    | 0.943     | 0.824     | 25       | 2        | 1        | 7        | –        |
|                           | STR strict | type 4 Outdated               | 53  | 0.962 | 0.898    | 0.975     | 0.923     | 39       | 2        | 0        | 12       | –        |
|                           | STR strict | type 5 Misinformation         | 55  | 0.945 | 0.891    | 0.947     | 0.943     | 27       | 0        | 3        | 25       | –        |
|                           | STR soft   | overall                       | 194 | 0.954 | 0.888    | 0.967     | 0.920     | 133      | 5        | 4        | 52       | –        |
|                           | STR soft   | type 1 No conflict            | 51  | 0.961 | 0.865    | 0.976     | 0.889     | 41       | 2        | 0        | 8        | –        |
|                           | STR soft   | type 2 Complementary          | 35  | 0.943 | 0.838    | 0.963     | 0.875     | 26       | 1        | 1        | 7        | –        |
|                           | STR soft   | type 4 Outdated               | 53  | 0.962 | 0.898    | 0.975     | 0.923     | 39       | 2        | 0        | 12       | –        |
|                           | STR soft   | type 5 Misinformation         | 55  | 0.945 | 0.891    | 0.947     | 0.943     | 27       | 0        | 3        | 25       | –        |
|                           | FG claim   | overall                       | 428 | 0.949 | 0.888    | 0.960     | 0.928     | 264      | 12       | 10       | 142      | –        |
|                           | FG claim   | type 1 No conflict            | 76  | 0.947 | 0.885    | 0.959     | 0.926     | 47       | 2        | 2        | 25       | –        |
|                           | FG claim   | type 2 Complementary          | 99  | 0.949 | 0.890    | 0.961     | 0.930     | 61       | 3        | 2        | 33       | –        |
|                           | FG claim   | type 3 Conflicting op.        | 102 | 0.961 | 0.912    | 0.971     | 0.941     | 66       | 3        | 1        | 32       | –        |
|                           | FG claim   | type 4 Outdated               | 68  | 0.897 | 0.762    | 0.925     | 0.837     | 43       | 3        | 4        | 18       | –        |
|                           | FG claim   | type 5 Misinformation         | 83  | 0.976 | 0.951    | 0.979     | 0.971     | 47       | 1        | 1        | 34       | –        |
| Human vs individual judge | Behaviour  | local/deepseek-r1-distill-32b | 650 | 0.735 | 0.339    | 0.820     | 0.503     | 391      | 41       | 131      | 87       | –        |
|                           | Behaviour  | local/mistral-small-4         | 650 | 0.643 | 0.229    | 0.720     | 0.506     | 299      | 133      | 99       | 119      | –        |
|                           | Behaviour  | local/qwen3.5-397b-a17b       | 649 | 0.743 | 0.405    | 0.812     | 0.592     | 361      | 70       | 97       | 121      | –        |

Table 22 gives all four binary comparison frames at every stratum, including the six reviewer pairs, the five conflict types, the per-model strata and the per-judge comparison, with positive- and negative-specific agreement and the full confusion cells. The reviewer-pair rows show that agreement is not uniform across annotators either: human behaviour agreement ranges from 0.819 to 1.000 across pairs, and the smallest pairs involve the reviewer who completed half an assigned queue, so their high coefficients rest on very few items and should not be read as evidence of superior reliability.

Table 23: **Complete continuous grounding-ratio agreement.** Per-sample Factual-Grounding ratios compared as continuous values. *exact* uses the configured tolerance rather than floating-point identity. Correlation measures association, not calibration, so it must be read with MAE/RMSE and the two means. Human–human rows are response-level over the 300-response frame; human–committee rows are judgment-level.

| Frame                  | Stratum                | $n$ | exact | MAE   | RMSE  | Pearson | Spearman | mean A | mean B |
|------------------------|------------------------|-----|-------|-------|-------|---------|----------|--------|--------|
| Human vs human         | overall                | 300 | 0.860 | 0.096 | 0.278 | 0.807   | 0.796    | 0.647  | 0.635  |
|                        | R1 / R2                | 83  | 0.855 | 0.074 | 0.224 | 0.887   | 0.871    | 0.601  | 0.676  |
|                        | R1 / R3                | 84  | 0.929 | 0.062 | 0.236 | 0.865   | 0.869    | 0.617  | 0.627  |
|                        | R1 / R4                | 21  | 0.905 | 0.048 | 0.163 | 0.937   | 0.923    | 0.647  | 0.694  |
|                        | R2 / R3                | 83  | 0.783 | 0.169 | 0.383 | 0.693   | 0.686    | 0.738  | 0.569  |
|                        | R2 / R4                | 19  | 1.000 | 0.000 | 0.000 | 1.000   | 1.000    | 0.697  | 0.697  |
|                        | R3 / R4                | 10  | 0.600 | 0.233 | 0.408 | 0.641   | 0.638    | 0.433  | 0.667  |
|                        | type 1 No conflict     | 63  | 0.984 | 0.016 | 0.126 | 0.965   | 0.965    | 0.640  | 0.656  |
|                        | type 2 Complementary   | 60  | 0.850 | 0.092 | 0.263 | 0.826   | 0.798    | 0.672  | 0.625  |
|                        | type 3 Conflicting op. | 58  | 0.793 | 0.135 | 0.329 | 0.696   | 0.680    | 0.645  | 0.654  |
|                        | type 4 Outdated        | 60  | 0.783 | 0.161 | 0.368 | 0.661   | 0.642    | 0.694  | 0.633  |
|                        | type 5 Misinformation  | 59  | 0.881 | 0.079 | 0.251 | 0.855   | 0.862    | 0.583  | 0.606  |
| Human vs committee     | overall                | 650 | 0.865 | 0.095 | 0.281 | 0.805   | 0.798    | 0.627  | 0.655  |
|                        | type 1 No conflict     | 133 | 0.932 | 0.050 | 0.207 | 0.901   | 0.885    | 0.657  | 0.662  |
|                        | type 2 Complementary   | 130 | 0.854 | 0.109 | 0.305 | 0.751   | 0.754    | 0.633  | 0.660  |
|                        | type 3 Conflicting op. | 128 | 0.797 | 0.118 | 0.294 | 0.765   | 0.752    | 0.604  | 0.654  |
|                        | type 4 Outdated        | 130 | 0.815 | 0.146 | 0.356 | 0.687   | 0.677    | 0.651  | 0.685  |
|                        | type 5 Misinformation  | 129 | 0.922 | 0.054 | 0.213 | 0.897   | 0.894    | 0.588  | 0.614  |
|                        | llama8b                | 328 | 0.845 | 0.109 | 0.299 | 0.766   | 0.755    | 0.675  | 0.713  |
|                        | qwen7b                 | 322 | 0.885 | 0.082 | 0.261 | 0.836   | 0.832    | 0.578  | 0.596  |
| Consensus vs committee | overall                | 258 | 0.938 | 0.045 | 0.195 | 0.909   | 0.906    | 0.648  | 0.636  |
|                        | type 1 No conflict     | 62  | 0.935 | 0.046 | 0.195 | 0.914   | 0.897    | 0.651  | 0.648  |
|                        | type 2 Complementary   | 51  | 0.922 | 0.069 | 0.252 | 0.836   | 0.844    | 0.650  | 0.641  |
|                        | type 3 Conflicting op. | 46  | 0.935 | 0.025 | 0.101 | 0.972   | 0.974    | 0.654  | 0.643  |
|                        | type 4 Outdated        | 47  | 0.894 | 0.089 | 0.280 | 0.816   | 0.805    | 0.681  | 0.642  |
|                        | type 5 Misinformation  | 52  | 1.000 | 0.000 | 0.000 | 1.000   | 1.000    | 0.606  | 0.606  |

Table 23 reports the continuous grounding-ratio comparisons. These are complementary to the claim-



level statistics rather than redundant: claim-level agreement asks whether reviewers make the same individual support decisions, while the ratio compares the proportion of an answer they judge grounded. A reviewer pair can agree on most claims yet differ on the ratio when their extracted claim sets differ in size.

Table 24: **Complete committee-internal agreement.** Agreement among the three judge models on the 300-response agreement frame. This is agreement among model judges under a fixed prompt and serving protocol; it is not human inter-annotator agreement and must not be reported as such.

| Subset             | Metric     | Judge pair                   | $n$ | agr   | $\kappa$        | $P_{pos}$ | $P_{neg}$ |
|--------------------|------------|------------------------------|-----|-------|-----------------|-----------|-----------|
| 300-response frame | Behaviour  | <i>alpha / mean pairwise</i> | 300 |       | 0.4568 / 0.4699 |           |           |
|                    |            | DeepSeek-32B vs Mistral-4    | 300 | 0.730 | 0.368           | 0.810     | 0.532     |
|                    |            | DeepSeek-32B vs Qwen-397B    | 300 | 0.853 | 0.611           | 0.903     | 0.703     |
|                    |            | Mistral-4 vs Qwen-397B       | 300 | 0.743 | 0.431           | 0.806     | 0.621     |
|                    | STR strict | <i>alpha / mean pairwise</i> | 215 |       | 0.6780 / 0.6874 |           |           |
|                    |            | DeepSeek-32B vs Mistral-4    | 215 | 0.809 | 0.613           | 0.839     | 0.766     |
|                    |            | DeepSeek-32B vs Qwen-397B    | 215 | 0.944 | 0.874           | 0.958     | 0.915     |
|                    |            | Mistral-4 vs Qwen-397B       | 215 | 0.791 | 0.575           | 0.825     | 0.740     |
|                    | FG claim   | <i>alpha / mean pairwise</i> | 477 |       | 0.8706 / 0.8705 |           |           |
|                    |            | DeepSeek-32B vs Mistral-4    | 477 | 0.950 | 0.886           | 0.963     | 0.924     |
|                    |            | DeepSeek-32B vs Qwen-397B    | 477 | 0.941 | 0.869           | 0.956     | 0.913     |
|                    |            | Mistral-4 vs Qwen-397B       | 477 | 0.937 | 0.857           | 0.953     | 0.903     |

Table 24 reports committee-internal agreement on the 300-response frame. The reliability ordering is clear: grounding is high ( $\alpha \approx 0.87$ ), recall intermediate ( $\approx 0.67$ ), behaviour lowest ( $\approx 0.46$ ). The behaviour disagreement is concentrated between specific judge pairs rather than spread evenly, which is consistent with the per-judge human comparison in Table 13.

## Q Complete Results Matrix

Refusal recall moves from 0.000–0.633 to 0.898–1.000, and several fine-tuned models correctly abstain on all 128 refusal-required items. Baseline Llama-3.1-8B under the minimal prompt shows why decision accuracy and refusal F1 must be read together: its decision accuracy of 0.825 looks respectable, but the model correctly refuses one of 128 refusal-required items, visible only in its refusal F1 of 0.015. The per-family Behaviour Adherence gains are +0.096 (Mistral-7B), +0.068 (Llama-3.1-8B), +0.047 (Qwen2.5-32B) and +0.000 (Qwen2.5-7B); the Qwen baselines already score 0.789 and 0.811 under the strict prompt against 0.660 for Mistral-7B, leaving less room for supervision to add.

Table 25 gives all twelve model×prompt cells, base and fine-tuned.

| Model        | Prompt  | Run  | Dec-Acc | Ref-F1 | BA    | FG    | STR   | AQ    | CATS-P | CATS-B |
|--------------|---------|------|---------|--------|-------|-------|-------|-------|--------|--------|
| Qwen2.5-7B   | strict  | base | 0.707   | 0.429  | 0.811 | 0.617 | 0.561 | 0.524 | 0.450  | 0.459  |
| Qwen2.5-7B   | strict  | SFT  | 0.967   | 0.911  | 0.780 | 0.838 | 0.687 | 0.677 | 0.630  | 0.684  |
| Qwen2.5-7B   | runtime | base | 0.784   | 0.361  | 0.768 | 0.651 | 0.648 | 0.541 | 0.453  | 0.375  |
| Qwen2.5-7B   | runtime | SFT  | 0.966   | 0.906  | 0.774 | 0.829 | 0.693 | 0.690 | 0.629  | 0.679  |
| Qwen2.5-7B   | minimal | base | 0.827   | 0.031  | 0.737 | 0.031 | 0.757 | 0.033 | 0.022  | 0.017  |
| Qwen2.5-7B   | minimal | SFT  | 0.958   | 0.881  | 0.768 | 0.837 | 0.686 | 0.680 | 0.617  | 0.659  |
| Qwen2.5-32B  | strict  | base | 0.855   | 0.367  | 0.789 | 0.851 | 0.733 | 0.724 | 0.574  | 0.432  |
| Qwen2.5-32B  | strict  | SFT  | 0.963   | 0.905  | 0.821 | 0.812 | 0.749 | 0.686 | 0.661  | 0.707  |
| Qwen2.5-32B  | runtime | base | 0.855   | 0.389  | 0.799 | 0.864 | 0.780 | 0.760 | 0.600  | 0.455  |
| Qwen2.5-32B  | runtime | SFT  | 0.962   | 0.901  | 0.804 | 0.814 | 0.758 | 0.701 | 0.658  | 0.716  |
| Qwen2.5-32B  | minimal | base | 0.826   | 0.015  | 0.747 | 0.108 | 0.773 | 0.105 | 0.086  | 0.057  |
| Qwen2.5-32B  | minimal | SFT  | 0.966   | 0.911  | 0.809 | 0.822 | 0.738 | 0.687 | 0.652  | 0.716  |
| Llama-3.1-8B | strict  | base | 0.762   | 0.255  | 0.721 | 0.705 | 0.634 | 0.577 | 0.426  | 0.363  |
| Llama-3.1-8B | strict  | SFT  | 0.973   | 0.926  | 0.805 | 0.852 | 0.727 | 0.706 | 0.680  | 0.744  |
| Llama-3.1-8B | runtime | base | 0.826   | 0.030  | 0.751 | 0.897 | 0.720 | 0.746 | 0.510  | 0.335  |
| Llama-3.1-8B | runtime | SFT  | 0.974   | 0.930  | 0.806 | 0.849 | 0.705 | 0.690 | 0.673  | 0.748  |
| Llama-3.1-8B | minimal | base | 0.825   | 0.015  | 0.743 | 0.074 | 0.710 | 0.055 | 0.029  | 0.018  |
| Llama-3.1-8B | minimal | SFT  | 0.976   | 0.934  | 0.795 | 0.857 | 0.714 | 0.694 | 0.672  | 0.732  |
| Mistral-7B   | strict  | base | 0.845   | 0.352  | 0.660 | 0.624 | 0.563 | 0.489 | 0.378  | 0.303  |
| Mistral-7B   | strict  | SFT  | 0.927   | 0.826  | 0.793 | 0.796 | 0.644 | 0.632 | 0.620  | 0.677  |
| Mistral-7B   | runtime | base | 0.826   | 0.200  | 0.676 | 0.739 | 0.590 | 0.566 | 0.398  | 0.299  |
| Mistral-7B   | runtime | SFT  | 0.929   | 0.831  | 0.760 | 0.762 | 0.613 | 0.593 | 0.587  | 0.657  |
| Mistral-7B   | minimal | base | 0.823   | 0.000  | 0.727 | 0.079 | 0.700 | 0.073 | 0.048  | 0.025  |
| Mistral-7B   | minimal | SFT  | 0.933   | 0.839  | 0.786 | 0.792 | 0.627 | 0.620 | 0.608  | 0.663  |

Table 25: **End-to-end results, all twelve model  $\times$  prompt cells** (24 rows). Dec-Acc is decision accuracy; GR-P/R/F1 are answer-positive; Ref-F1 is refusal-positive.  $n = 736$  throughout. Oracle-condition rows are summarised in Table 19.